

MODEL EVALUATION

How Do We Evaluate?

- what is overall "accuracy"
- are we better predictiong (green or red) labels?
- how much better are we compared with random ("coin" flipping)

Many Metrics




figure reprinted from www.kdnuggets.com with explicit permission of the editor

A Numerical Dataset

| object x_i | Height (H) | Weight (W) | Foot (F) | Label (L) |
|-----------------|---------------|---------------|-------------|--------------|
| x_1 | 5.00 | 100 | 6 | green |
| x_2 | 5.50 | 150 | 8 | green |
| x_3 | 5.33 | 130 | 7 | green |
| x_4 | 5.75 | 150 | 9 | green |
| x_5 | 6.00 | 180 | 13 | red |
| x_6 | 5.92 | 190 | 11 | red |
| x_7 | 5.58 | 170 | 12 | red |
| x_8 | 5.92 | 165 | 10 | red |

Code for the Dataset


```
import pandas as pd

data = pd.DataFrame(
    {"id": [ 1,2,3,4,5,6,7,8] ,
     "Label": ["green", "green", "green", "green",
               "red", "red", "red", "red"] ,
     "Height": [5,5.5,5.33,5.75,6.00,5.92,5.58,5.92] ,
     "Weight": [100,150,130,150,180,190,170,165] ,
     "Foot": [6, 8, 7, 9, 13, 11, 12, 10]} ,
    columns=["id", "Height", "Weight",
             "Foot", "Label"])
```

```
>> data
```

| | id | Height | Weight | Foot | Label |
|---|----|--------|--------|------|-------|
| 0 | 1 | 5.00 | 100 | 6 | green |
| 1 | 2 | 5.50 | 150 | 8 | green |
| 2 | 3 | 5.33 | 130 | 7 | green |
| 3 | 4 | 5.75 | 150 | 9 | green |
| 4 | 5 | 6.00 | 180 | 13 | red |
| 5 | 6 | 5.92 | 190 | 11 | red |
| 6 | 7 | 5.58 | 170 | 12 | red |
| 7 | 8 | 5.92 | 165 | 10 | red |

A Dataset Illustration



Three Models

- objects:

$$x_1, x_2, x_3, x_4, x_5, x_6, x_7, x_8$$


- "green" is 1, 'red" is 0
- true labels: [1,1,1,1,0,0,0,0]
- want to compare 3 models:
 1. Model 1: [0,0,0,1,0,0,1,1]
 2. Model 2: [0,0,1,1,0,1,1,1]
 3. Model 3: [1,1,1,1,1,1,1,0]

Python Code

```
import pandas as pd
data = pd.DataFrame(
{"id": [ 1,2,3,4,5,6,7,8] ,
 "Label": ["green", "green", "green", "green",
           "red", "red", "red", "red"] ,
 "Height": [5,5.5,5.33,5.75,6.00,5.92,5.58,5.92] ,
 "Weight": [100,150,130,150,180,190,170,165] ,
 "Foot": [6, 8, 7, 9, 13, 11, 12, 10]} ,
 columns=["id", "Height", "Weight",
          "Foot", "Label"])
data["Class"] = data["Label"].apply(lambda x: \
                                   1 if x=="green" else 0)
y_true = data["Class"].values
# assume that we got predictions from 3 models:
y_pred_1 = [0,0,0,1,0,0,1,1]
y_pred_2 = [0,0,1,1,0,1,1,1]
y_pred_3 = [1,1,1,1,1,1,1,0]

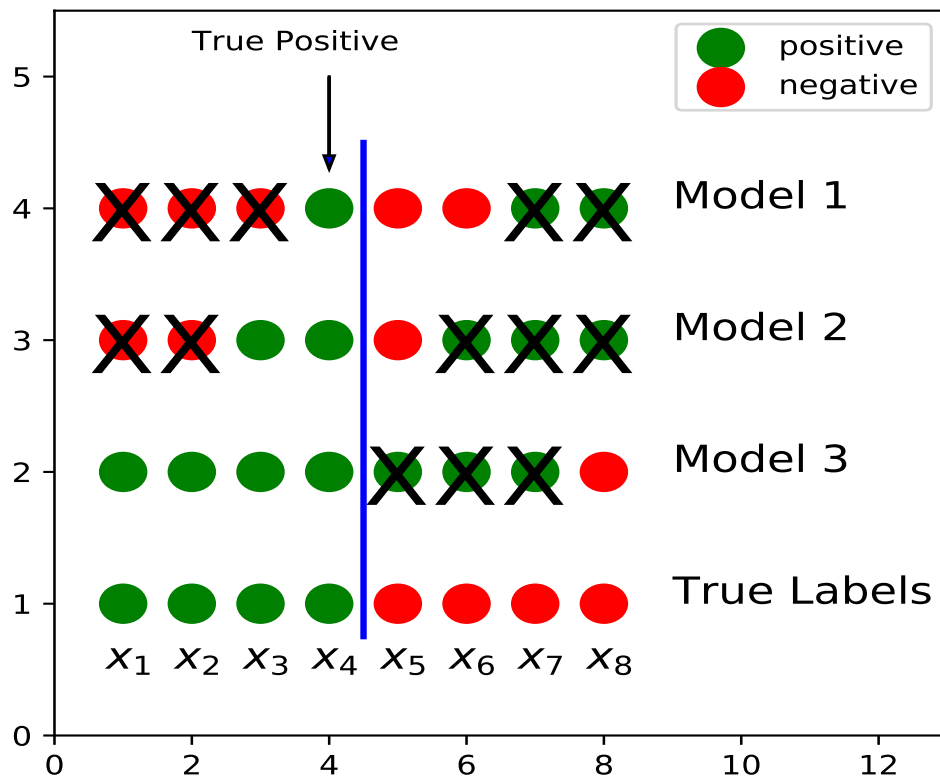
>> y_true
[1, 1, 1, 1, 0, 0, 0, 0]
>> y_pred_1
[0, 0, 0, 1, 0, 0, 1, 1]
```


Comparison of Models



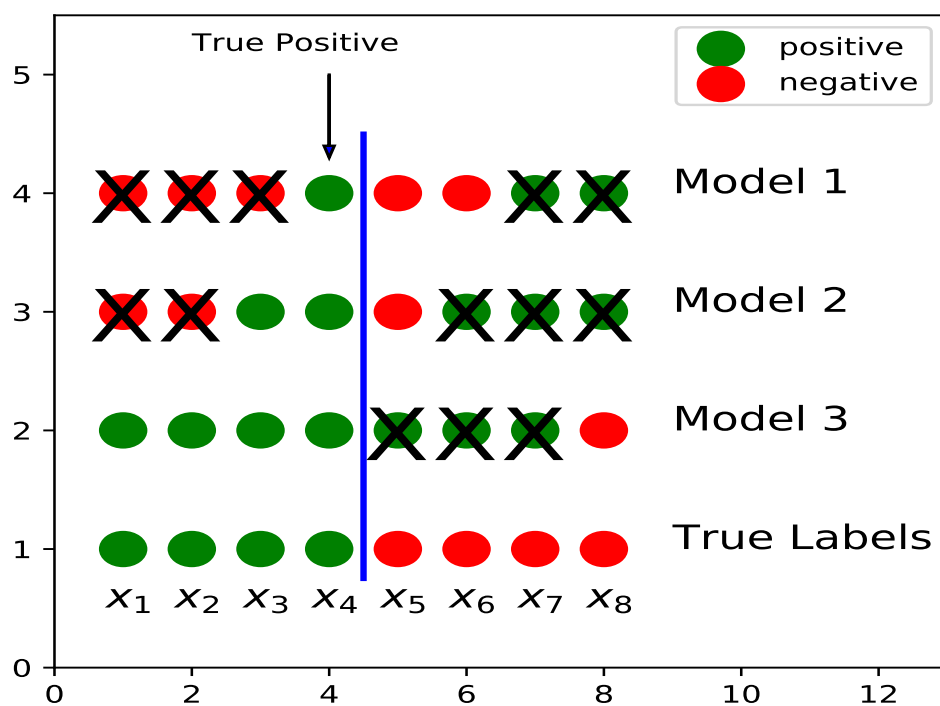
- "same" accuracy for 1 and 2
- sometimes 3 is the "best"
- how do we compare?

True Positives (TP)



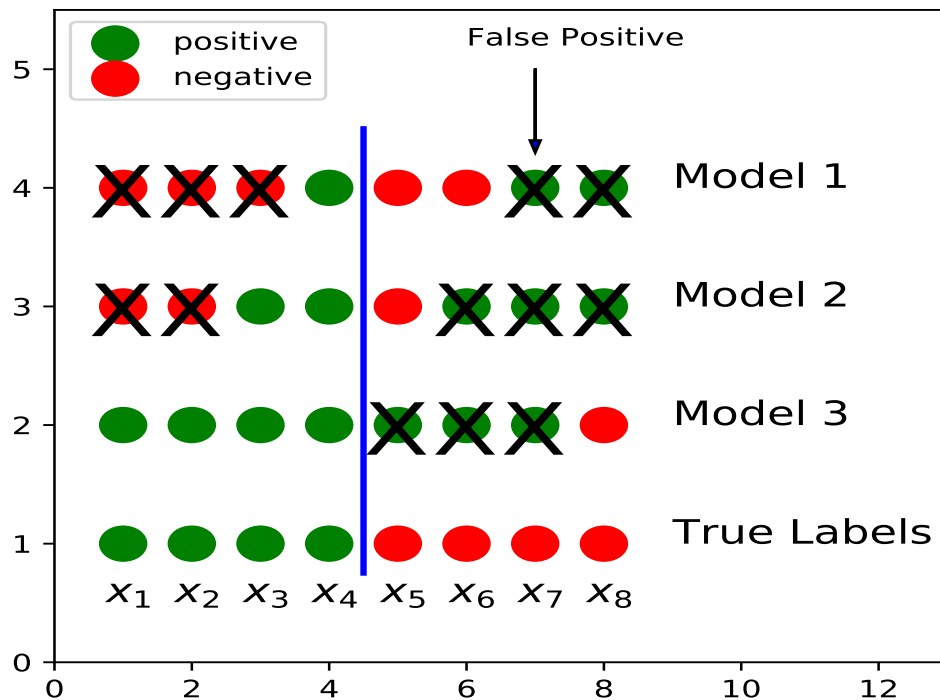
- predicted positive and true labels are positive

True Positives (TP)



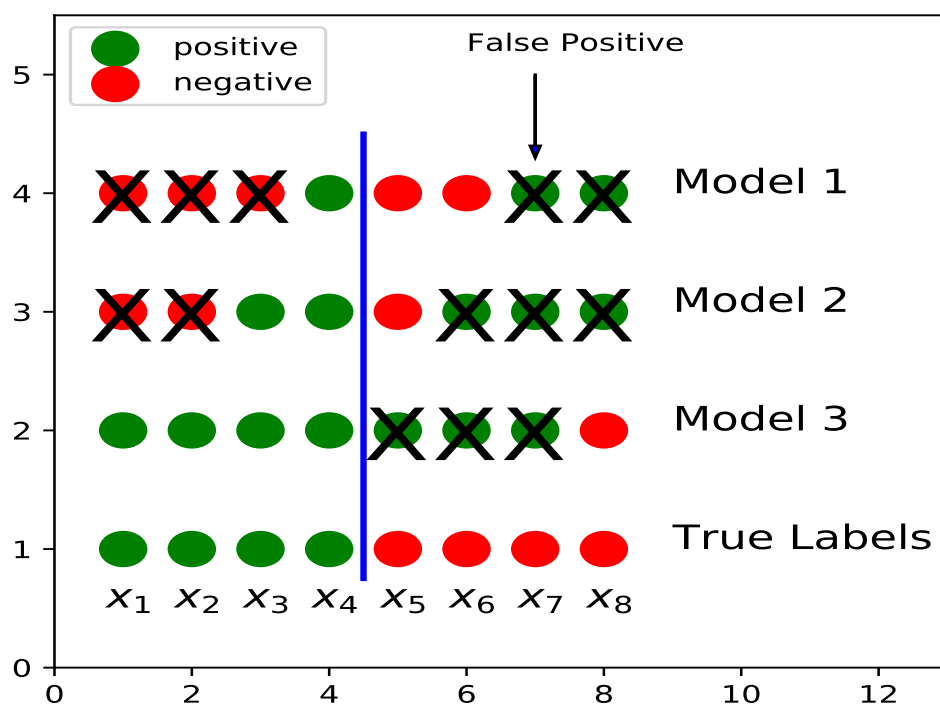
| Model | True Positive | Actual Positive |
|-------|----------------------|----------------------|
| 1 | x_4 | x_1, x_2, x_3, x_4 |
| 2 | x_3, x_4 | x_1, x_2, x_3, x_4 |
| 3 | x_1, x_2, x_3, x_4 | x_1, x_2, x_3, x_4 |

False Positives (FP)



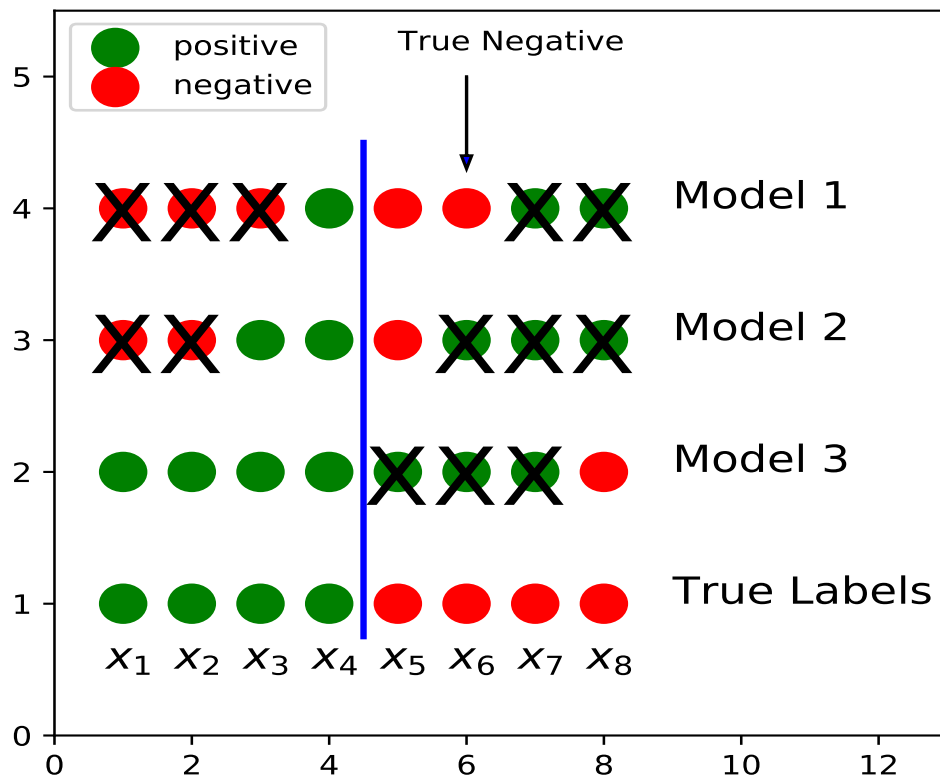
- predicted positive but true labels are negative
- "Type I" error ("false alarm")

False Positives (FP)



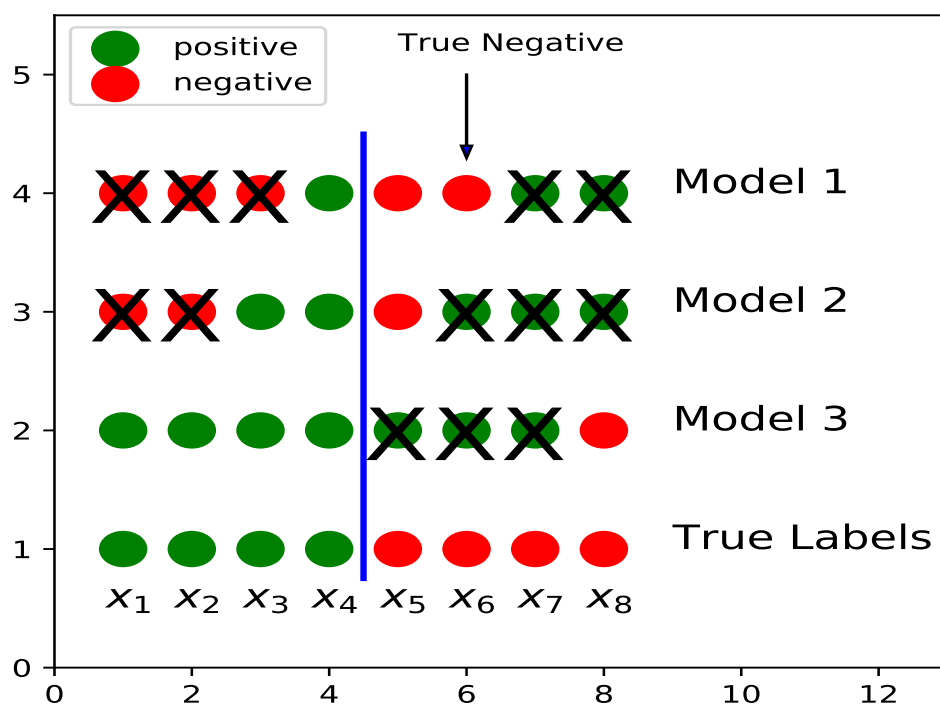
| Model | False Positive | Actual Negative |
|-------|-----------------|----------------------|
| 1 | x_7, x_8 | x_5, x_6, x_7, x_8 |
| 2 | x_6, x_7, x_8 | x_5, x_6, x_7, x_8 |
| 3 | x_5, x_6, x_7 | x_5, x_6, x_7, x_8 |

True Negative (TN)



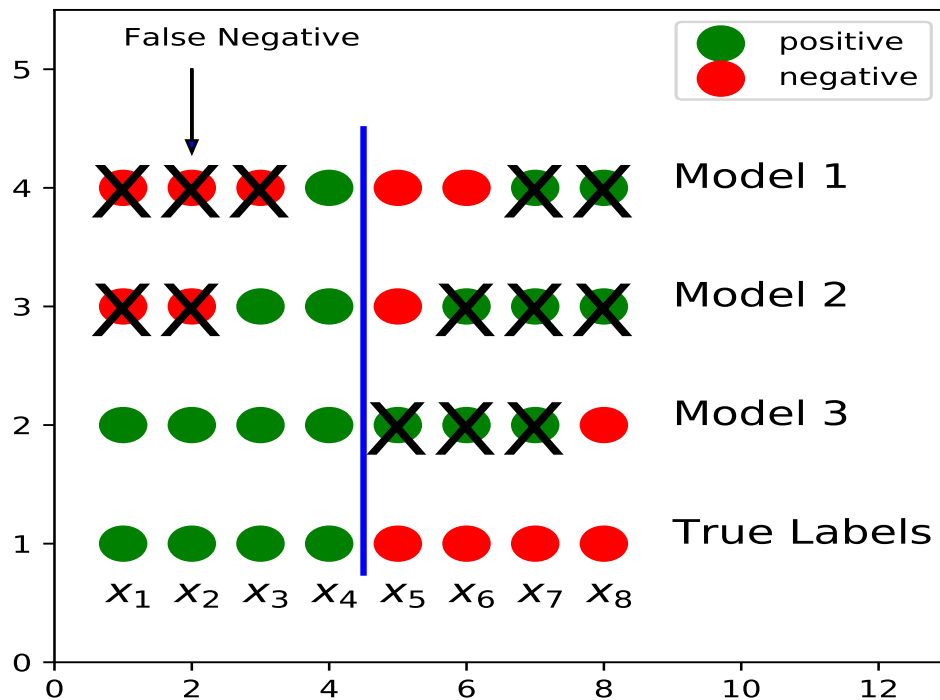
- predicted negative and true labels are negative

True Negatives (TN)



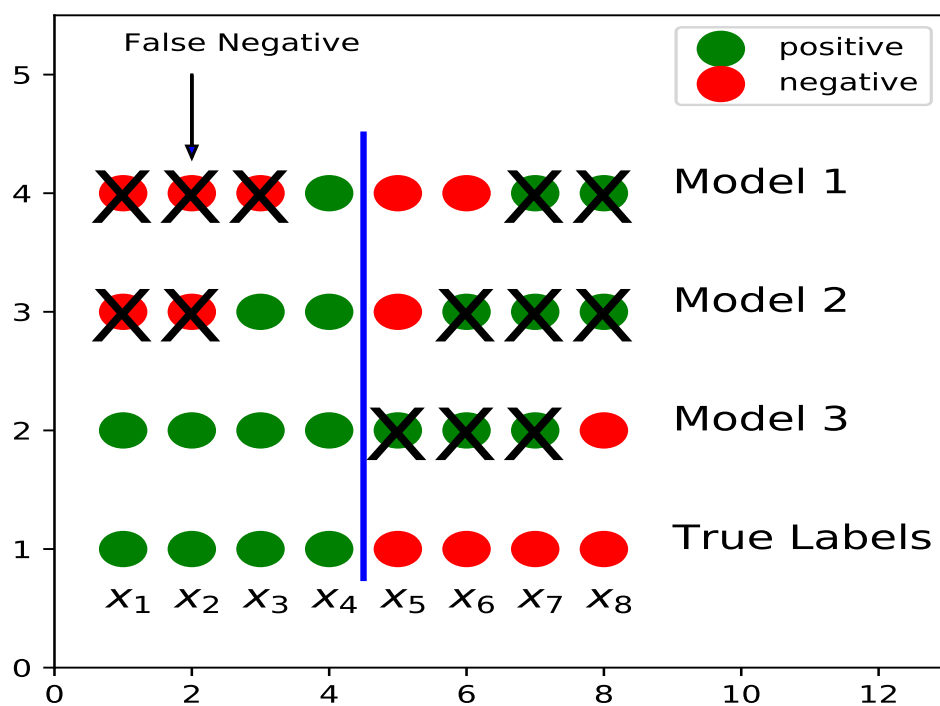
| Model | True Negative | Actual Negative |
|-------|---------------|----------------------|
| 1 | x_5, x_6 | x_5, x_6, x_7, x_8 |
| 2 | x_5 | x_5, x_6, x_7, x_8 |
| 3 | x_8 | x_5, x_6, x_7, x_8 |

False Negative (FN)



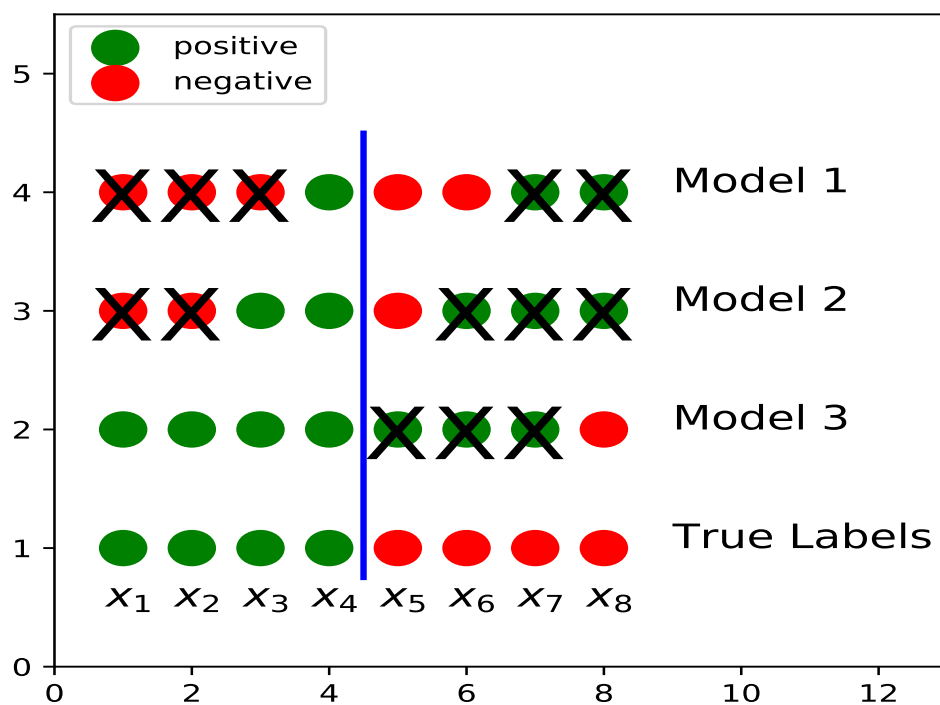
- predicted negative, but true labels are positive
- "Type II" error ("miss")

False Negative (FN)



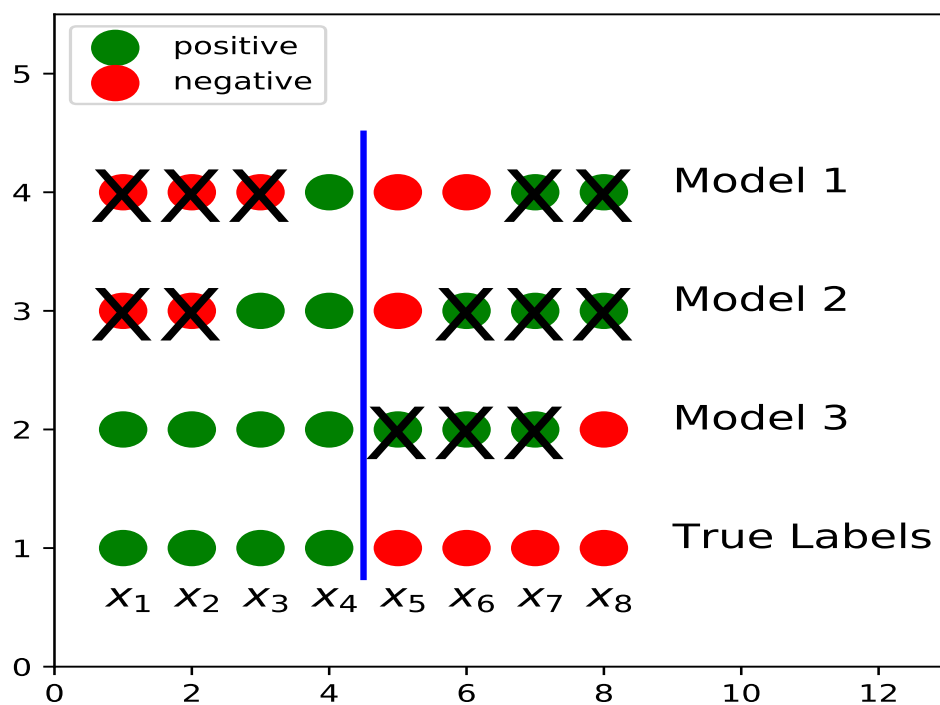
| Model | False Negative | Actual Positive |
|-------|-----------------|----------------------|
| 1 | x_1, x_2, x_3 | x_1, x_2, x_3, x_4 |
| 2 | x_1, x_2 | x_1, x_2, x_3, x_4 |
| 3 | none | x_1, x_2, x_3, x_4 |

Pos. & Neg. Summary



| Model | TP | FP | TN | FN |
|-------|----------------------|-----------------|------------|-----------------|
| 1 | x_4 | x_7, x_8 | x_5, x_6 | x_1, x_2, x_3 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | x_5 | x_1, x_2 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | x_8 | none |

Pos. & Neg. Summary



| Model | TP | FP | TN | FN |
|-------|----------------------|-----------------|------------|-----------------|
| 1 | x_4 | x_7, x_8 | x_5, x_6 | x_1, x_2, x_3 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | x_5 | x_1, x_2 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | x_8 | none |

Confusion Matrix

- each row represents predictions
- each column represents actual class
- $C = [C_{ij}]$
- C_{ij} - observations in group j predicted for group i

Visual Representation

- rows represent predictions
- columns represent actual class

$$C = \begin{bmatrix} \text{TN} & \text{FP} \\ \text{FN} & \text{TP} \end{bmatrix}$$

TN/()

- total positive $P = TP + FN$
- total negative $N = TN + FP$

Confusion Matrices

| Model | TP | FP | TN | FN |
|-------|----------------------|-----------------|------------|-----------------|
| 1 | x_4 | x_7, x_8 | x_5, x_6 | x_1, x_2, x_3 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | x_5 | x_1, x_2 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | x_8 | none |

$$C = \begin{bmatrix} \text{TN} & \text{FP} \\ \text{FN} & \text{TP} \end{bmatrix}$$

$$C_1 = \begin{bmatrix} 2 & 2 \\ 3 & 1 \end{bmatrix}, \quad C_2 = \begin{bmatrix} 1 & 3 \\ 2 & 2 \end{bmatrix}, \quad C_3 = \begin{bmatrix} 1 & 3 \\ 0 & 4 \end{bmatrix}$$

Python Code

```
import pandas as pd
from sklearn.metrics import confusion_matrix

data = pd.DataFrame(
    {'id': [ 1,2,3,4,5,6,7,8],
     'Label': ['green','green','green','green',
               'red','red','red','red'],
     'Height': [5, 5.5, 5.33, 5.75,
                6.00, 5.92, 5.58, 5.92],
     'Weight': [100, 150, 130, 150,
                180, 190, 170, 165],
     'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
    columns = ['id', 'Height', 'Weight',
               'Foot', 'Label'] )

data['Class'] = data['Label'].apply(lambda x: 1
                                    if x=='green' else 0)

y_true = data['Class'].values
# assume that we got predictions from 3 models:
y_pred_1 = [0,0,0,1,0,0,1,1]
y_pred_2 = [0,0,1,1,0,1,1,1]
y_pred_3 = [1,1,1,1,1,1,1,0]
cf_1 = confusion_matrix(y_true,y_pred_3)
cf_2 = confusion_matrix(y_true,y_pred_3)
cf_3 = confusion_matrix(y_true,y_pred_3)
```

```
>> cf_3
```

```
[[1 3]
```

```
 [0 4]]
```

True Positive Rate (TPR)

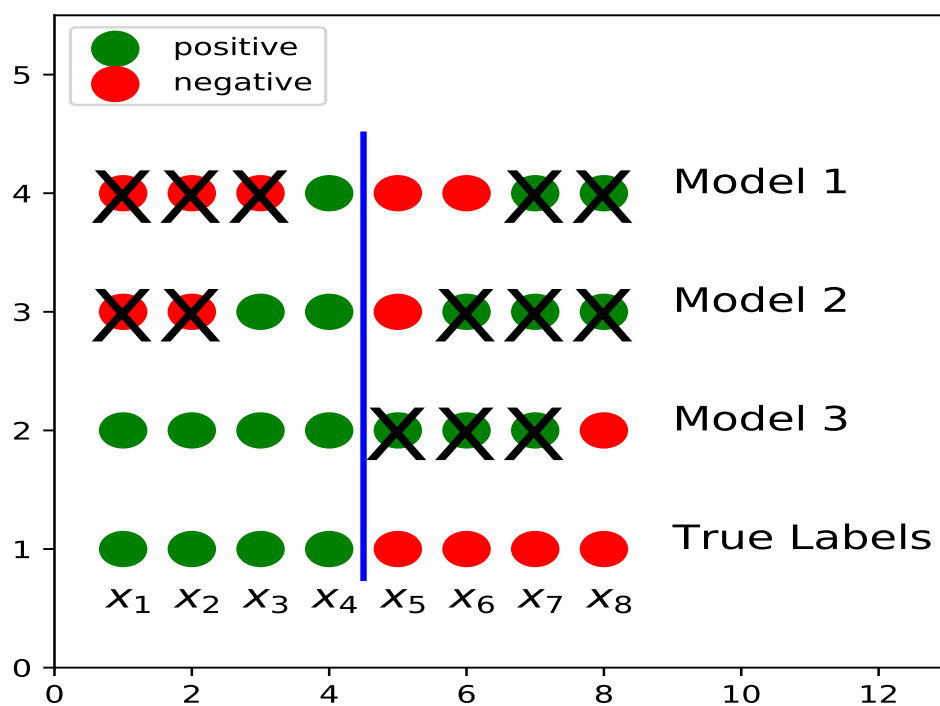
- sensitivity, recall, or hit rate

$$\text{TPR} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

| Model | TP | FN | TPR |
|-------|----------------------|-----------------|------|
| 1 | x_4 | x_1, x_2, x_3 | 0.25 |
| 2 | x_3, x_4 | x_1, x_2 | 0.5 |
| 3 | x_1, x_2, x_3, x_4 | none | 1 |

- fraction of positive labels predicted correctly

$$\text{TPR} = \text{TP} / (\text{TP} + \text{FN})$$



| Model | TP | FN | TPR |
|-------|----------------------|-----------------|------|
| 1 | x_4 | x_1, x_2, x_3 | 0.25 |
| 2 | x_3, x_4 | x_1, x_2 | 0.5 |
| 3 | x_1, x_2, x_3, x_4 | none | 1 |

Python Code

```
import pandas as pd
from sklearn.metrics import recall_score

data = pd.DataFrame(
    {'id': [ 1,2,3,4,5,6,7,8],
     'Label': ['green','green','green','green',
               'red','red','red','red'],
     'Height': [5, 5.5, 5.33, 5.75,
                6.00, 5.92, 5.58, 5.92],
     'Weight': [100, 150, 130, 150,
                180, 190, 170, 165],
     'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
    columns = ['id', 'Height', 'Weight',
               'Foot', 'Label'] )

data['Class'] = data['Label'].apply(lambda x: 1
                                     if x=='green' else 0)

y_true = data['Class'].values
# assume that we got predictions from 3 models:
y_pred_1 = [0,0,0,1,0,0,1,1]
y_pred_2 = [0,0,1,1,0,1,1,1]
y_pred_3 = [1,1,1,1,1,1,1,0]

tpr_1 = recall_score(y_true, y_pred_1)
tpr_2 = recall_score(y_true, y_pred_2)
tpr_3 = recall_score(y_true, y_pred_3)

>> print(tpr_1, tpr_2, tpr_3)

0.25 0.5 1.0
```

True Negative Rate (TNR)

- ”specificity” or ”selectivity”

$$\text{TNR} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

| Model | FP | TN | TNR |
|-------|-----------------|------------|------|
| 1 | x_7, x_8 | x_5, x_6 | 0.50 |
| 2 | x_6, x_7, x_8 | x_5 | 0.25 |
| 3 | x_5, x_6, x_7 | x_8 | 0.25 |

- fraction of negative labels predicted correctly


Positive Predicted Value (PPV)

- ”precision”

$$\text{PPV} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

| Model | TP | FP | PPV |
|-------|----------------------|-----------------|------|
| 1 | x_4 | x_7, x_8 | 0.33 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | 0.40 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | 0.57 |

$$PPV = TP / (TP + FP)$$



| Model | TP | FP | PPV |
|-------|----------------------|-----------------|------|
| 1 | x_4 | x_7, x_8 | 0.33 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | 0.40 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | 0.57 |

Python Code

```
import pandas as pd
from sklearn.metrics import precision_score

data = pd.DataFrame(
    {'id': [ 1,2,3,4,5,6,7,8],
     'Label': ['green','green','green','green',
               'red','red','red','red'],
     'Height': [5, 5.5, 5.33, 5.75,
                6.00, 5.92, 5.58, 5.92],
     'Weight': [100, 150, 130, 150,
                180, 190, 170, 165],
     'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
    columns = ['id', 'Height', 'Weight',
               'Foot', 'Label'] )

data['Class'] = data['Label'].apply(lambda x: 1
                                    if x=='green' else 0)

y_true = data['Class'].values
# assume that we got predictions from 3 models:
y_pred_1 = [0,0,0,1,0,0,1,1]
y_pred_2 = [0,0,1,1,0,1,1,1]
y_pred_3 = [1,1,1,1,1,1,1,0]

ppv_1 = precision_score(y_true, y_pred_1)
ppv_2 = precision_score(y_true, y_pred_2)
ppv_3 = precision_score(y_true, y_pred_3)
```

```
>> print(ppv_1, ppv_2, ppv_3)
```

```
0.33333333333333 0.4 0.571428571429
```

Negative Predicted Value (NPV)

$$\text{NPV} = \frac{\text{TN}}{\text{TN} + \text{FN}}$$

| Model | TN | FN | NPV |
|-------|------------|-----------------|------|
| 1 | x_5, x_6 | x_1, x_2, x_3 | 0.40 |
| 2 | x_5 | x_1, x_2 | 0.33 |
| 3 | x_8 | none | 1.0 |

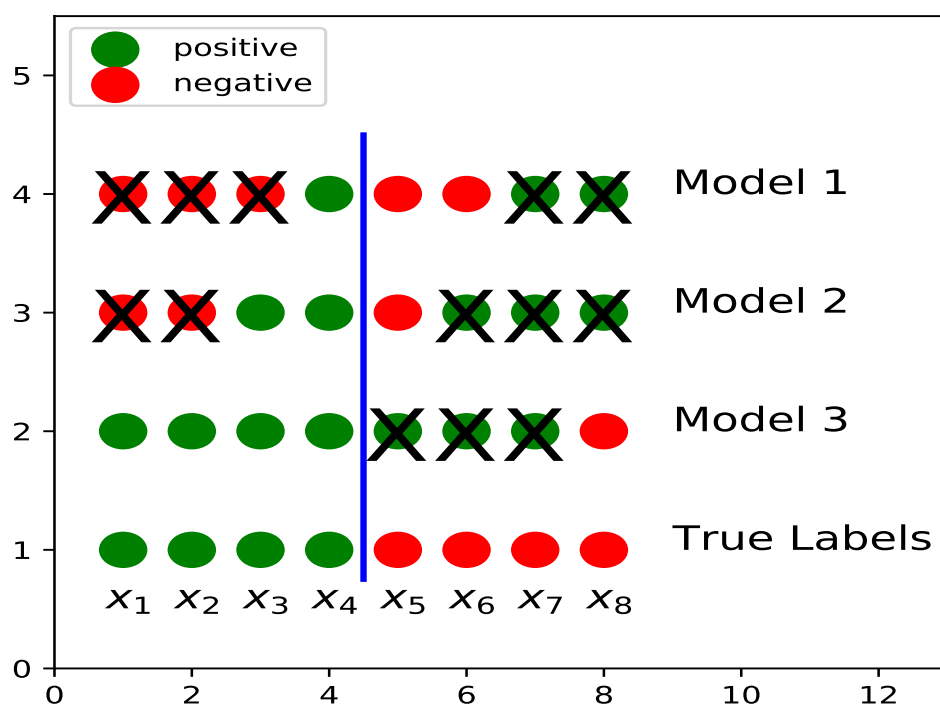
Accuracy

$$\text{ACC} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

| Model | TP | FP | TN | FN | ACC |
|-------|----------------------|-----------------|------------|-----------------|-------|
| 1 | x_4 | x_7, x_8 | x_5, x_6 | x_1, x_2, x_3 | 0.375 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | x_5 | x_1, x_2 | 0.375 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | x_8 | none | 0.625 |

- fraction of all labels predicted correctly
- models 1 and 2 - same accuracy, different precision

$$\text{ACC} = (\text{TP} + \text{TN}) / \text{ALL}$$



| Model | TP | FP | TN | FN | ACC |
|-------|----------------------|-----------------|------------|-----------------|-------|
| 1 | x_4 | x_7, x_8 | x_5, x_6 | x_1, x_2, x_3 | 0.375 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | x_5 | x_1, x_2 | 0.375 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | x_8 | none | 0.625 |

Python Code

```
import pandas as pd
from sklearn.metrics import accuracy_score

data = pd.DataFrame(
    {'id': [ 1,2,3,4,5,6,7,8],
     'Label': ['green','green','green','green',
               'red','red','red','red'],
     'Height': [5, 5.5, 5.33, 5.75,
                6.00, 5.92, 5.58, 5.92],
     'Weight': [100, 150, 130, 150,
                180, 190, 170, 165],
     'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
    columns = ['id', 'Height', 'Weight',
               'Foot', 'Label'] )

data['Class'] = data['Label'].apply(lambda x: 1
                                    if x=='green' else 0)

y_true = data['Class'].values
# assume that we got predictions from 3 models:
y_pred_1 = [0,0,0,1,0,0,1,1]
y_pred_2 = [0,0,1,1,0,1,1,1]
y_pred_3 = [1,1,1,1,1,1,1,0]

acc_1 = accuracy_score(y_true, y_pred_1)
acc_2 = accuracy_score(y_true, y_pred_2)
acc_3 = accuracy_score(y_true, y_pred_3)

>> print(acc_1, acc_2, acc_3)

0.375 0.375 0.625
```

F_1 Score

- harmonic mean of precision and sensitivity

$$F_1 = 2 \cdot \frac{\text{PPV} \cdot \text{TPR}}{\text{PPV} + \text{TPR}}$$

$$= \frac{2 \cdot \text{TP}}{2 \cdot \text{TP} + \text{FP} + \text{FN}}$$

| Model | TP | FP | FN | F_1 |
|-------|----------------------|-----------------|-----------------|-------|
| 1 | x_4 | x_7, x_8 | x_1, x_2, x_3 | 0.29 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | x_1, x_2 | 0.44 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | none | 0.73 |

Python Code

```
import pandas as pd
from sklearn.metrics import f1_score

data = pd.DataFrame(
    {'id': [ 1,2,3,4,5,6,7,8],
     'Label': ['green','green','green','green',
               'red','red','red','red'],
     'Height': [5, 5.5, 5.33, 5.75,
                6.00, 5.92, 5.58, 5.92],
     'Weight': [100, 150, 130, 150,
                180, 190, 170, 165],
     'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
    columns = ['id', 'Height', 'Weight',
               'Foot', 'Label'] )

data['Class'] = data['Label'].apply(lambda x: 1
                                    if x=='green' else 0)

y_true = data['Class'].values
# assume that we got predictions from 3 models:
y_pred_1 = [0,0,0,1,0,0,1,1]
y_pred_2 = [0,0,1,1,0,1,1,1]
y_pred_3 = [1,1,1,1,1,1,1,0]

f1_1 = f1_score(y_true,y_pred_3)
f1_2 = f1_score(y_true,y_pred_3)
f1_3 = f1_score(y_true,y_pred_3)

>> print(f1_1, f1_2, f1_3)

0.2857142857 0.4444444444 0.7272727272
```

Comparing Models

| Metric | Model 1 | Model 2 | Model 3 |
|-------------------|---------|---------|---------|
| recall (TPR) | 0.25 | 0.5 | 1 |
| specificity (TPR) | 0.5 | 0.25 | 0.25 |
| precision (PPV) | 0.33 | 0.4 | 0.57 |
| accuracy | 0.375 | 0.375 | 0.625 |
| F_1 | 0.29 | 0.44 | 0.73 |

- choice of model depends on the metric

Additional Measures

- False Negative Rate:

$$\text{FNR} = 1 - \text{TPR}$$

- False Positive Rate:

$$\text{FPR} = 1 - \text{TNR}$$

- False Discovery Rate:

$$\text{FDR} = 1 - \text{PPV}$$

- False Omission Rate:

$$\text{FOR} = 1 - \text{NPV}$$

False Negative Rate (FNR)

- ”miss” rate

$$\text{FNR} = \frac{\text{FN}}{\text{FN} + \text{TP}}$$

| Model | TP | FN | FNR |
|-------|----------------------|-----------------|------|
| 1 | x_4 | x_1, x_2, x_3 | 0.75 |
| 2 | x_3, x_4 | x_1, x_2 | 0.5 |
| 3 | x_1, x_2, x_3, x_4 | none | 0 |

False Positive Rate (FPR)

- ”fall-out”

$$\text{FPR} = \frac{\text{FP}}{\text{FP} + \text{TN}}$$

| Model | FP | TN | FPR |
|-------|-----------------|------------|------|
| 1 | x_7, x_8 | x_5, x_6 | 0.5 |
| 2 | x_6, x_7, x_8 | x_5 | 0.75 |
| 3 | x_5, x_6, x_7 | x_8 | 0.75 |

False Discovery Rate (FDR)

$$\text{FDR} = \frac{\text{FP}}{\text{FP} + \text{TP}}$$

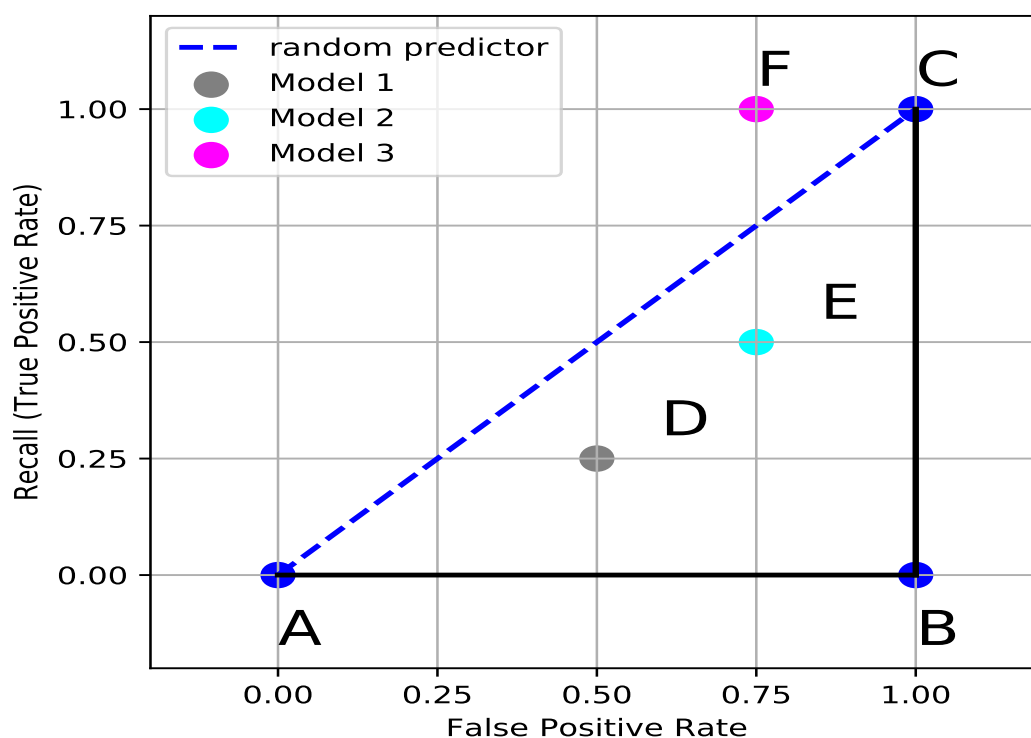
| Model | TP | FP | FDR |
|-------|----------------------|-----------------|------|
| 1 | x_4 | x_7, x_8 | 0.67 |
| 2 | x_3, x_4 | x_6, x_7, x_8 | 0.60 |
| 3 | x_1, x_2, x_3, x_4 | x_5, x_6, x_7 | 0.43 |

False Omission Rate (FOR)

$$\text{FOR} = \frac{\text{FN}}{\text{FN} + \text{TN}}$$

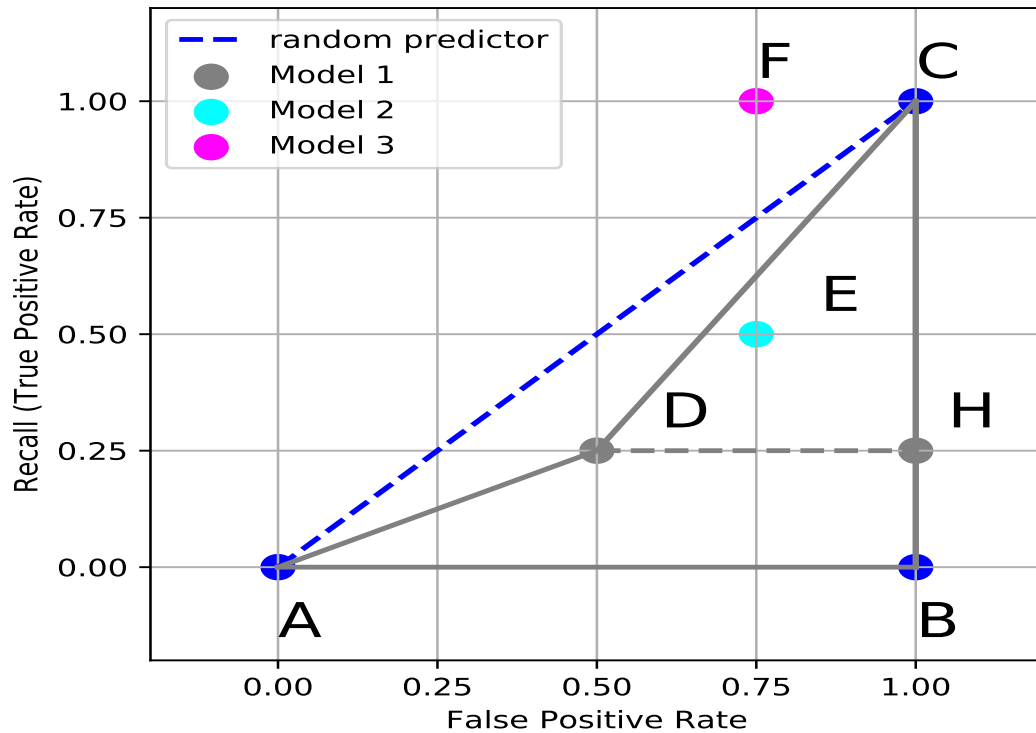
| Model | TN | FN | FOR |
|-------|------------|-----------------|------|
| 1 | x_5, x_6 | x_1, x_2, x_3 | 0.60 |
| 2 | x_5 | x_1, x_2 | 0.66 |
| 3 | x_8 | none | 0 |

ROC/AUC Curve



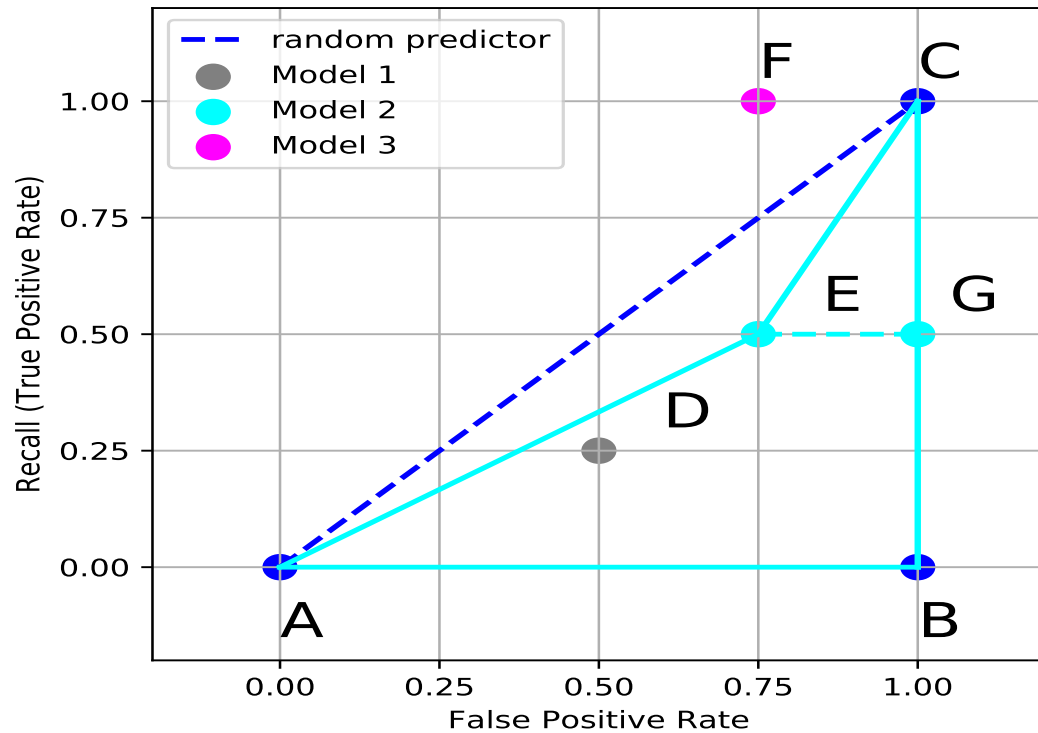
- receiver operating characteristic (ROC) - describe binary classifiers
- area under curve (AUC) - compare classifiers vs. random

AUC For Model 1



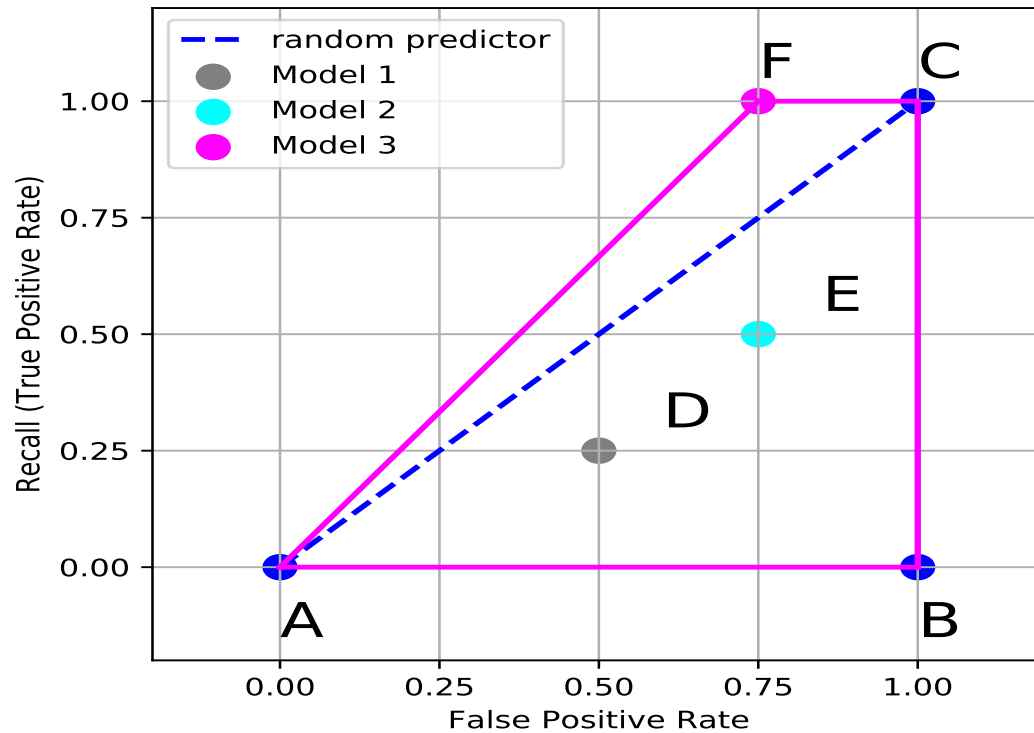
$$\begin{aligned}
 \text{area}(ABCD) &= \text{area}(ABDH) + \text{area}(DHC) \\
 &= \frac{(1 + 0.5)}{2} \cdot 0.25 + 0.5 \cdot 0.5 \cdot 0.75 \\
 &= 0.1875 + 0.1875 = 0.375
 \end{aligned}$$

AUC For Model 2



$$\begin{aligned}
 \text{area}(ABCE) &= \text{area}(ABGE) + \text{area}(GCE) \\
 &= \frac{(1 + 0.25)}{2} \cdot 0.5 + 0.25 \cdot 0.5 \cdot 0.5 \\
 &= 0.3125 + 0.0625 = 0.375
 \end{aligned}$$

AUC For Model 3



$$\text{area}(ABCF) = \frac{(1 + 0.25)}{2} = 0.625$$

Python Code

```
import pandas as pd
from sklearn.metrics import roc_auc_score

data = pd.DataFrame(
    {'id': [ 1,2,3,4,5,6,7,8],
     'Label': ['green','green','green','green',
               'red','red','red','red'],
     'Height': [5, 5.5, 5.33, 5.75,
                6.00, 5.92, 5.58, 5.92],
     'Weight': [100, 150, 130, 150,
                180, 190, 170, 165],
     'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
    columns = ['id', 'Height', 'Weight',
               'Foot', 'Label'] )

data['Class'] = data['Label'].apply(lambda x: 1
                                    if x=='green' else 0)

y_true = data['Class'].values
# assume that we got predictions from 3 models:
y_pred_1 = [0,0,0,1,0,0,1,1]
y_pred_2 = [0,0,1,1,0,1,1,1]
y_pred_3 = [1,1,1,1,1,1,1,0]

auc_1 = roc_auc_score(y_true, y_pred_1)
auc_2 = roc_auc_score(y_true, y_pred_2)
auc_3 = roc_auc_score(y_true, y_pred_3)

>> print(auc_1, auc_2, auc_3)

0.375  0.375  0.625
```

Concepts Check:

- (a) true and false positive
- (b) true and false negatives
- (c) sensitivity (or recall)
- (d) specificity, precision
- (e) type I and II error
- (f) confusion matrix
- (g) F_1 score
- (h) receiver operating characteristic (ROC)
- (i) area under curve (AUC)