

NAIVE BAYESIAN CLASSIFICATION

Bayes Formula

$$\underbrace{P(C|x)} = \underbrace{P(x|C)} \cdot \underbrace{P(C)} / \underbrace{P(x)}$$

posterior = likelihood · prior / evidence

- one of main formulae in probability and statistics
- used for classification for labels
- Naive Bayesian classifier

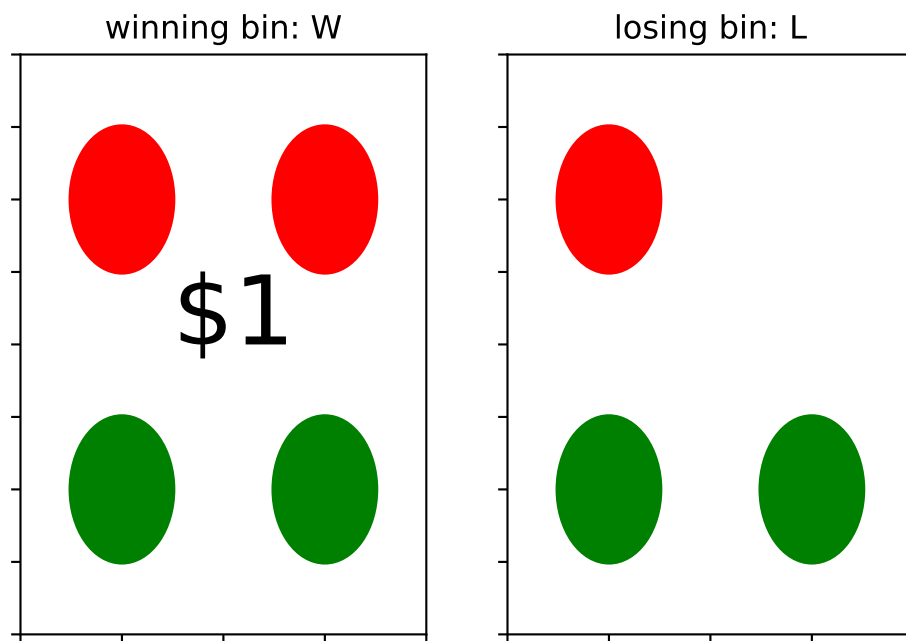
Bayes Formula

$$\underbrace{P(C|x)} = \underbrace{P(x|C)} \cdot \underbrace{P(C)} / \underbrace{P(x)}$$

posterior = likelihood · prior / evidence


- $P(C|x)$ - conditional probability of class C given training inputs x
- $P(x|C)$ - conditional probability of training inputs x given class C
- $P(C)$ - probability of class C
- $P(x)$ - probability of the training data x

Bayes Formula Example



- two bins: W and L
- W: 2 red, 2 green tokens, \$1
- L: 1 red, 2 green tokens


Bayes Formula Example



- equally likely choices
- prior probabilities:


$$P(W) = P(L) = 0.5$$

Bayes Formula Example



- can see one token from a bin
- token is 'green'
- what is $P(W | g)$?


Intuition



- L has a higher percentage of green tokens
- likelihoods:


$$P(g|W) = 0.5, \quad P(g|L) = 2/3$$

Intuition



- a green token means more likely that we are looking at bin L


Applying Formula



$$\underbrace{P(W|g)} = \underbrace{P(g|W)} \cdot \underbrace{P(W)} / \underbrace{P(g)}$$

posterior = likelihood · prior / evidence

Applying Formula



$$\begin{aligned}
 P(W|g) &= \frac{P(g|W) \cdot P(W)}{P(g|W)P(W) + P(g|L) \cdot P(L)} \\
 &= \frac{0.5 \cdot 0.5}{0.5 \cdot 0.5 + (2/3) \cdot 0.5} = 0.43
 \end{aligned}$$

Example: "Spam" Mail

- 40% of email is "spam" (s)
- 30% of spam contains "free" keyword (f)
- 1% of no-spam ($\neg s$) contains f keyword
- a new email contains f
- what is $P(s|f)$ - probability it is a spam?

Example: "Spam" Mail (cont'd)

$$P(S) = 0.4$$

$$P(\neg S) = 0.6$$

$$P(f|s) = 0.3$$

$$P(f|\neg S) = 0.01$$

$$P(s|f) = ?$$

- prior probability $p(s) = 0.4$
- intuition: $p(s|f) > 0.4$
- why? - we saw f

Computing $P(S|f)$

$$\begin{aligned} P(S|f) &= \frac{P(f|S) \cdot P(S)}{P(f)} \\ &= \frac{P(f|S) \cdot P(S)}{P(f|S) \cdot P(S) + P(f|\neg S) \cdot P(\neg S)} \\ &= \frac{0.3 \cdot 0.4}{0.3 \cdot 0.4 + 0.01 \cdot 0.6} \\ &= \frac{0.12}{0.12 + 0.06} \\ &= 2/3 \end{aligned}$$

- note $P(S|f) > P(S)$

Bayes Formula Example

Day	Weather	Play
1	sunny	no
2	rainy	yes
3	sunny	yes
4	rainy	no
5	sunny	yes
6	overcast	yes
7	sunny	yes
8	overcast	yes
9	rainy	no
10	rainy	yes

- assume Weather = sunny
- what is $P(\text{yes} \mid \text{sunny})$?

Python Code

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import \
    LabelEncoder
from sklearn.naive_bayes import \
    MultinomialNB

df = pd.DataFrame(
    {"Day": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
     "Weather": ["sunny", "rainy",
                  "sunny", "rainy",
                  "sunny", "overcast",
                  "sunny", "overcast",
                  "rainy", "rainy"],
     "Play": ["no", "yes", "yes",
               "no", "yes", "yes",
               "yes", "yes",
               "no", "yes"]},
    columns=["Day", "Weather",
             "Play"] )
```

Computing Priors

Day	Weather	Play
1	sunny	no
2	rainy	yes
3	sunny	yes
4	rainy	no
5	sunny	yes
6	overcast	yes
7	sunny	yes
8	overcast	yes
9	rainy	no
10	rainy	yes

$$P(\text{yes}) = 7/10$$

$$P(\text{no}) = 3/10$$

- unconditional probabilities

Computing Priors in Python

```
>> df_play = df["Play"].value_counts()
>> df_play
yes      7
no       3
Name: Play, dtype: int64
```

$$P(\text{yes}) = 7/10$$

$$P(\text{no}) = 3/10$$

Computing Evidence

Day	Weather	Play
1	sunny	no
2	rainy	yes
3	sunny	yes
4	rainy	no
5	sunny	yes
6	overcast	yes
7	sunny	yes
8	overcast	yes
9	rainy	no
10	rainy	yes

$$P(\text{sunny}) = 4/10$$

$$P(\text{rainy}) = 4/10$$

$$P(\text{overcast}) = 2/10$$

Computing Evidence in Python

```
>> df_weather = df.groupby(["Weather",  
                             "Play"]).size()
```

```
>> df_weather
```

Weather	Play	
overcast	yes	2
rainy	no	2
	yes	2
sunny	no	1
	yes	3

```
dtype: int64
```

$$P(\text{sunny}) = 4/10$$

$$P(\text{rainy}) = 4/10$$

$$P(\text{overcast}) = 2/10$$

- unconditional probabilities

Computing Likelihoods

Day	Weather	Play
1	sunny	no
2	rainy	yes
3	sunny	yes
4	rainy	no
5	sunny	yes
6	overcast	yes
7	sunny	yes
8	overcast	yes
9	rainy	no
10	rainy	yes

$$P(\text{sunny} \mid \text{yes}) = 3/7$$

- conditional probabilities

Computing Likelihood in Python

```
>> df_weather = df.groupby(["Weather",  
                             "Play"]).size()
```

```
>> df_weather
```

Weather	Play	
overcast	yes	2
rainy	no	2
	yes	2
sunny	no	1
	yes	3

```
dtype: int64
```

$$P(\text{sunny} \mid \text{yes}) = 3/7$$

Applying the Formula

- likelihood: $P(\text{sunny} | \text{yes}) = \frac{3}{7}$
- prior: $P(\text{yes}) = \frac{7}{10}$
- evidence: $P(\text{sunny}) = \frac{4}{10}$

$$\begin{aligned} P(\text{yes} | \text{sunny}) &= \frac{P(\text{sunny} | \text{yes})P(\text{yes})}{P(\text{sunny})} \\ &= \frac{3}{7} \cdot \frac{7}{10} / \frac{4}{10} \\ &= \frac{3}{4} \end{aligned}$$

Naive Bayesian Classification

- assume multiple features x_1, \dots, x_n with class label C
- want to compute $P(C|x_1, \dots, x_n)$
- how? use the Bayes formula:

$$P(C|x_1, \dots, x_n) = \frac{P(x_1, \dots, x_n|C) \cdot P(C)}{P(x)}$$
$$\propto P(x_1, \dots, x_n|C) \cdot P(C)$$

- Q: where is difficulty?
- A: joint probability $P(x_1, \dots, x_n|C)$

Naive Bayesian Classification (cont'd)

- assume feature independence

$$P(x_1, \dots, x_n | C) = P(x_1 | C) \cdots P(x_n | C)$$

- Bayes simplifies to

$$P(C | x_1, \dots, x_n) \propto [P(x_1 | C) \cdots P(x_n | C)] P(C)$$

- assign $x^* = (a_1, \dots, a_n)$ to class C^* with maximum value $P(C^* | x^*)$

$$C^* = \operatorname{argmax}_C [P(a_1 | C) \cdots P(a_n | C)] P(C)$$

Naive Bayesian in SkLearn

- three classifiers:
 - (a) *GaussianNB*: features are continuous and follow normal distribution
 - (b) *MultinomialNB*: features are frequencies
 - (c) *BernoulliNB*: features are boolean

Naive Bayesian Example

- additional features: wind and temperature

Day	Weather	Temperature	Wind	Play
1	sunny	hot	low	no
2	rainy	mild	high	yes
3	sunny	cold	low	yes
4	rainy	cold	high	no
5	sunny	cold	high	yes
6	overcast	mild	low	yes
7	sunny	hot	low	yes
8	overcast	hot	high	yes
9	rainy	hot	high	no
10	rainy	mild	low	yes

- $x^* = (\text{sunny}, \text{cold}, \text{low})$?

Python Code

```
import pandas as pd
from sklearn.preprocessing import \
    LabelEncoder
from sklearn.naive_bayes import \
    MultinomialNB

df = pd.DataFrame(
{"Day": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
"Weather": ["sunny", "rainy", "sunny",
            "rainy", "sunny", "overcast", "sunny",
            "overcast", "rainy", "rainy"],
"Temperature": ["hot", "mild", "cold", "cold",
                "cold", "mild", "hot", "hot", "hot", "mild"],
"Wind": ["low", "high", "low", "high", "high",
         "low", "low", "high", "high", "low"],
"Play": ["no", "yes", "yes", "no", "yes",
         "yes", "yes", "yes", "no", "yes"]},
columns=["Day", "Weather", "Temperature",
        "Wind", "Play"] )
```

• $x^* = (\text{sunny}, \text{cold}, \text{low})?$

Naive Bayesian Example (cont'd)

- $x^* = (\text{sunny}, \text{cold}, \text{low})$?
- need a label C^* to maximize $P(\text{sunny}|C^*) \cdot P(\text{cold}|C^*) \cdot P(\text{low}|C^*)$
- need conditional probabilities
 1. weather: $P(\text{sunny}|C)$
 2. temperature: $P(\text{cold}|C)$
 3. wind: $P(\text{low}|C)$
- compute during learning phase

Cond. Probability: Weather

Weather	Play = yes	Play = no
sunny	3/7	1/3
overcast	2/7	0
rainy	2/7	2/3

$$P(\text{sunny} \mid \text{yes}) = 3/7$$

$$P(\text{sunny} \mid \text{no}) = 1/3$$

Cond. Probability: Weather (Python)

```
>> df_weather = data.groupby(["Weather",  
                              "Play"]).size()
```

```
>> df_weather
```

Weather	Play	
overcast	yes	2
rainy	no	2
	yes	2
sunny	no	1
	yes	3

```
dtype: int64
```

$$P(\text{sunny} \mid \text{yes}) = 3/7$$

$$P(\text{sunny} \mid \text{no}) = 1/3$$

Cond. Probability: Temperature

Temperature	Play = yes	Play = no
cold	2/7	1/3
mild	3/7	0
hot	2/7	2/3

$$P(\text{cold} \mid \text{yes}) = 2/7$$

$$P(\text{cold} \mid \text{no}) = 1/3$$

Cond. Probability: Temperature (Python)

```
>> df_temp = data.groupby(["Temperature",  
                           "Play"]).size()
```

```
>> df_temp
```

Temperature	Play	
cold	no	1
	yes	2
hot	no	2
	yes	2
mild	yes	3

```
dtype: int64
```

$$P(\text{cold} \mid \text{yes}) = 2/7$$

$$P(\text{cold} \mid \text{no}) = 1/3$$

Cond. Probability: Wind

Wind	Play = yes	Play = no
high	3/7	2/3
low	4/7	1/3

$$P(\text{low} \mid \text{yes}) = 4/7$$

$$P(\text{low} \mid \text{no}) = 1/3$$

Cond. Probability: Wind (Python)

```
>> df_wind = data.groupby(["Wind",  
                           "Play"]).size()
```

```
>> df_wind
```

```
Wind  Play
```

```
high  no      2
```

```
      yes     3
```

```
low   no      1
```

```
      yes     4
```

```
dtype: int64
```

$$P(\text{low} \mid \text{yes}) = 4/7$$

$$P(\text{low} \mid \text{no}) = 1/3$$

Computing Label

- $x^* = (\text{sunny}, \text{cold}, \text{low})?$

$$\begin{aligned} P(\text{yes}|x^*) &= P(\text{sunny}|\text{yes})P(\text{cold}|\text{yes}) \\ &\quad \cdot P(\text{low}|\text{yes})P(\text{yes}) \\ &= 3/7 \cdot 2/7 \cdot 4/7 \cdot 0.7 = 0.049 \end{aligned}$$

$$\begin{aligned} P(\text{no}|x^*) &= P(\text{sunny}|\text{no})P(\text{cold}|\text{no}) \\ &\quad \cdot P(\text{low}|\text{no})P(\text{no}) \\ &= 1/3 \cdot 1/3 \cdot 1/3 \cdot 0.3 = 0.011 \end{aligned}$$

- $P(\text{yes}|x^*) > P(\text{no}|x^*)$
- therefore, $x^* \mapsto \text{yes}$

Python Code

```
import numpy as np
import pandas as pd
from sklearn.naive_bayes import MultinomialNB
from sklearn.preprocessing import LabelEncoder

data = pd.DataFrame(
    {'Day': [1,2,3,4,5,6,7,8,9,10],
     'Weather': ['sunny','rainy','sunny','rainy',
                 'sunny','overcast','sunny','overcast',
                 'rainy','rainy'],
     'Temperature': ['hot', 'mild', 'cold','cold','cold',
                    'mild','hot','hot', 'hot','mild'],
     'Wind': ['low','high','low','high','high',
             'low','low', 'high','high','low'],
     'Play': ['no', 'yes','yes','no','yes',
             'yes','yes','yes','no','yes']},
    columns = ['Day','Weather','Temperature','Wind','Play'])

input_data = data[['Weather', 'Temperature', 'Wind']]
dummies = [pd.get_dummies(data[c]) for c in input_data.columns]
binary_data = pd.concat(dummies, axis=1)
X = binary_data[0:10].values
le = LabelEncoder()
Y = le.fit_transform(data['Play'].values)
NB_classifier = MultinomialNB().fit(X, Y)
new_instance = np.asmatrix([0,0,1,1,0,0,0,1])
prediction = NB_classifier.predict(new_instance)
```

```
ipdb> prediction[0]
```


```
1
```

Naive Bayesian: Continuous Case

object x_i	Height (H)	Weight (W)	Foot (F)	Label (L)
x_1	5.00	100	6	green
x_2	5.50	150	8	green
x_3	5.33	130	7	green
x_4	5.75	150	9	green
x_5	6.00	180	13	red
x_6	5.92	190	11	red
x_7	5.58	170	12	red
x_8	5.92	165	10	red

- $P(\text{green}) = P(\text{red}) = 0.5$
- $(H=6, W=160, F=10) \mapsto ?$

Naive Bayesian: Continuous Case (cont'd)



$(H=6, W=160, F=10) \mapsto ?$

Computing the Label

- assume H, W, F independent

$$\begin{aligned} P(C | x^*) &= \frac{P(x^* | C) P(C)}{P(x^*)} \\ &= \frac{P(H | C) P(W | C) P(F | C) P(C)}{P(x^*)} \end{aligned}$$

- term $P(x^*)$ same for each class
- assign C^* to x^* if for $C \neq C^*$:

$$\begin{aligned} &P(H | C^*) P(W | C^*) P(F | C^*) P(C^*) \\ &> \\ &P(H | C) P(W | C) P(F | C) P(C) \end{aligned}$$

Computing the Parameters

- assume Gaussian

label C	height H		weight W		foot F	
	μ	σ	μ	σ	μ	σ
green	5.39	0.27	132.5	20.46	7.5	1.12
red	5.86	0.16	176.25	9.60	11.5	1.12

Python Code for Probabilities

```
import pandas as pd

data = pd.DataFrame( {'id': [ 1,2,3,4,5,6,7,8],
    'Label': ['green', 'green', 'green', 'green',
              'red', 'red', 'red', 'red'],
    'Height': [5, 5.5, 5.33, 5.75,
               6.00, 5.92, 5.58, 5.92],
    'Weight': [100, 150, 130, 150,
               180, 190, 170, 165],
    'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
    columns = ['id', 'Height', 'Weight',
               'Foot', 'Label'] )

df_1 = data[data['Label']=='green']
h_mu_1 = df_1['Height'].values.mean()
h_sigma_1 = df_1['Height'].values.std()
w_mu_1 = df_1['Weight'].values.mean()
w_sigma_1 = df_1['Weight'].values.std()
f_mu_1 = df_1['Foot'].values.mean()
f_sigma_1 = df_1['Foot'].values.std()

df_2 = data[data['Label']=='red']
h_mu_2 = df_2['Height'].values.mean()
h_sigma_2 = df_2['Height'].values.std()
w_mu_2 = df_2['Weight'].values.mean()
w_sigma_2 = df_2['Weight'].values.std()
f_mu_2 = df_2['Foot'].values.mean()
f_sigma_2 = df_2['Foot'].values.std()
```

Python Code for Cond. Probabilities

```
from scipy.stats import norm

h, w, f = 6, 160, 10
p_green = 0.5; p_red = 0.5

prob_h_green = norm.pdf((h - h_mu_1)/h_sigma_1)
prob_w_green = norm.pdf((w - w_mu_1)/w_sigma_1)
prob_f_green = norm.pdf((f - f_mu_1)/f_sigma_1)

prob_h_red = norm.pdf((h - h_mu_2)/h_sigma_2)
prob_w_red = norm.pdf((w - w_mu_2)/w_sigma_2)
prob_f_red = norm.pdf((f - f_mu_2)/f_sigma_2)

# unnormalized probabilities
posterior_red = p_red * prob_h_red * prob_w_red * prob_f_red
posterior_green = p_green * prob_h_green * prob_w_green * prob_f_green

normalized_red = posterior_red / (posterior_red + posterior_green)
normalized_green = posterior_green / (posterior_red + posterior_green)

ipdb> normalized_red
0.96
```

$$x^* = (H = 6, W = 160, F = 10) \\ \mapsto \text{red}$$

Python Code for Naive Bayesian

```
import numpy as np
import pandas as pd
from sklearn.naive_bayes import GaussianNB

data = pd.DataFrame( {'id': [ 1,2,3,4,5,6,7,8],
                      'Label': ['green', 'green', 'green', 'green',
                                'red', 'red', 'red', 'red'],
                      'Height': [5, 5.5, 5.33, 5.75,
                                6.00, 5.92, 5.58, 5.92],
                      'Weight': [100, 150, 130, 150,
                                180, 190, 170, 165],
                      'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
                      columns = ['id', 'Height', 'Weight',
                                'Foot', 'Label'] )

X = data[['Height', 'Weight', 'Foot']].values
Y = data[['Label']].values

NB_classifier = GaussianNB().fit(X, Y)
new_instance = np.asmatrix(np.asmatrix([6, 160, 10]))
prediction = NB_classifier.predict(new_instance)

ipdb> prediction[0]
red
```

Naive Bayesian: IRIS

```
import pandas as pd
import numpy as np
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

url = r'https://archive.ics.uci.edu/ml/' + \
      r'machine-learning-databases/iris/iris.data'

iris_feature_names = ['sepal-length', 'sepal-width',
                      'petal-length', 'petal-width']
data = pd.read_csv(url, names=['sepal-length', 'sepal-width',
                              'petal-length', 'petal-width', 'Class'])

class_labels = ['Iris-versicolor', 'Iris-virginica']
data = data[data['Class'].isin(class_labels)]

X = data[iris_feature_names].values
le = LabelEncoder()
Y = le.fit_transform(data['Class'].values)

X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
                                                    test_size=0.5, random_state=3)
NB_classifier = GaussianNB().fit(X_train, Y_train)
prediction = NB_classifier.predict(X_test)
error_rate = np.mean(prediction != Y_test)

ipdb> error_rate
0.04
```

Concepts Check:

- (a) conditional probability
- (b) Bayes formula
- (c) prior and posterior probabilities
- (d) feature independence assumption
- (e) naive bayesian classification
- (f) discrete cases
- (g) continuous case