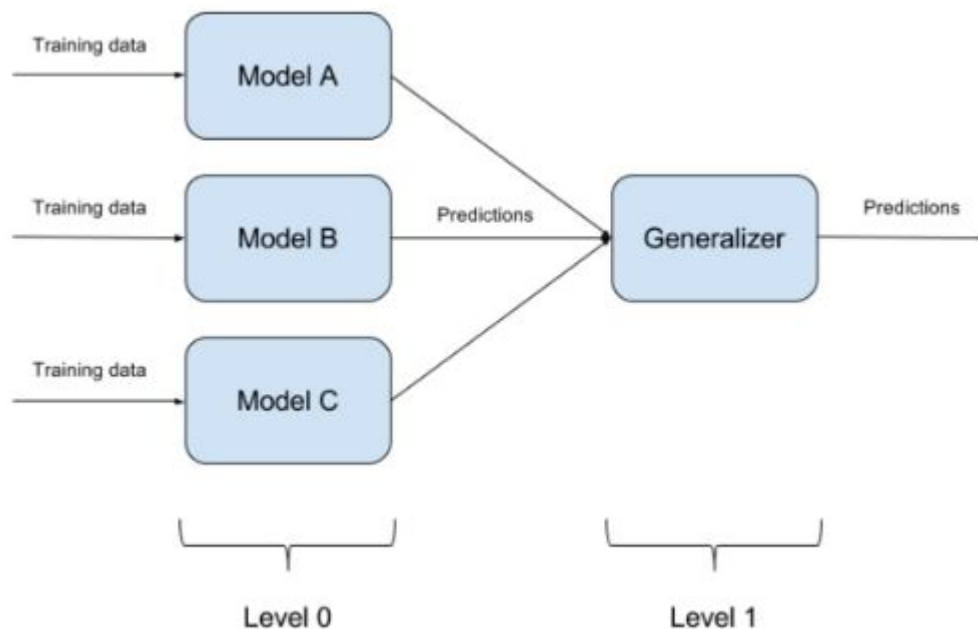


ENSEMBLE CLASSIFIERS

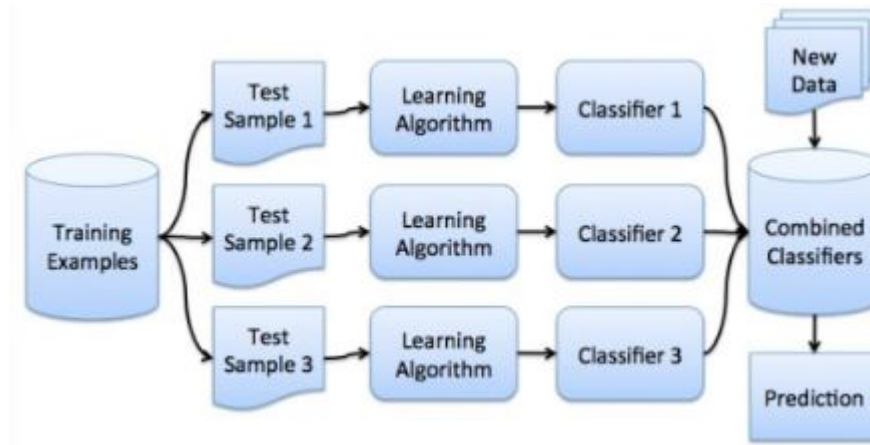
Ensemble Methods



- generate ("base learners")
- combine learners
- goal: increase accuracy

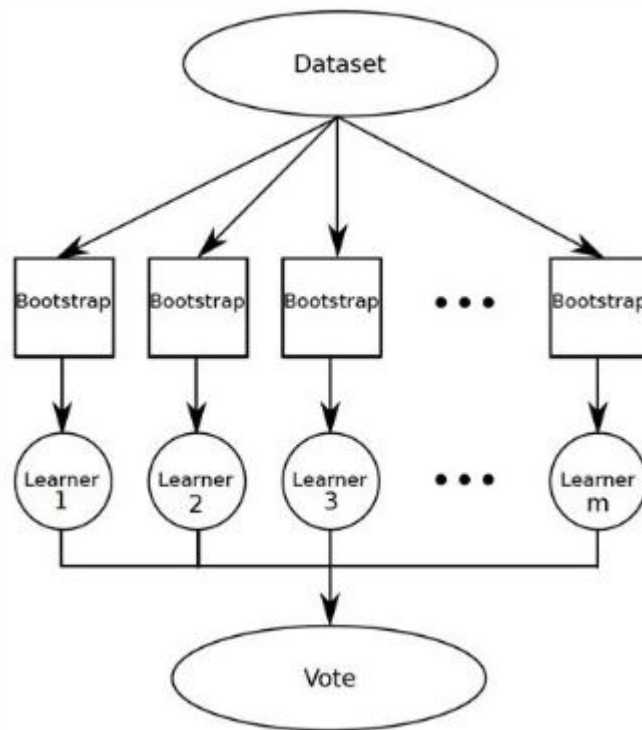
figure reprinted from www.kdnuggets.com with explicit permission of the editor

Sequential Ensemble



- learners are generated sequentially (Adaboost)

Parallel Ensemble



- learners are generated in parallel (random forest)

figure reprinted from www.kdnuggets.com with explicit permission of the editor

Random Forests vs. Decision Trees


Decision Trees:

- inputs and outputs
- want classification for labels
- decision tree classifier

Random Forests

- build partial decision (sub)trees
- decide by ensemble of (sub)trees
- conceptually "similar" to kNN

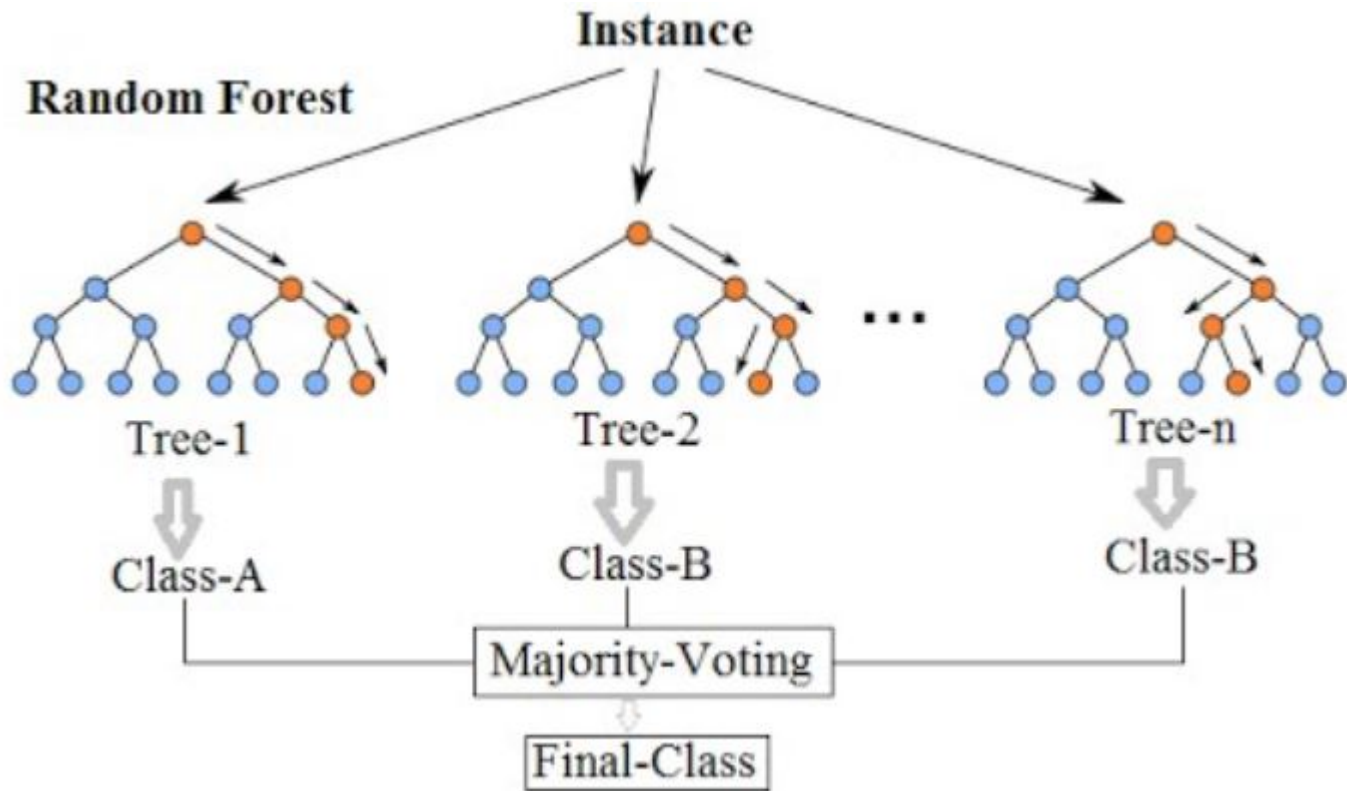
Random Forest (RF)



- a random sampling of data is used for each tree
- random subsets of features in splitting

figure reprinted from www.kdnuggets.com with explicit permission of the editor

Training & Predictions



- randomly sample subsets (“bagging”)
- fit models and predict

figure reprinted from www.kdnuggets.com with explicit permission of the editor

Bagging vs. Boosting

- *bagging*:

- (a) samples with replacements
- (b) equal weights of models
- (c) majority voting
- (d) ex: Random Forest

- *boosting*:

- (a) new models from older ones
- (b) weighted voting
- (c) ex: AdaBoost and LogBoost

Advantages/Disadvantages

- advantages:
 - (a) lower variance
 - (b) better predictions
- disadvantages:
 - (a) do not train well on small datasets
 - (b) results hard to interpret
 - (c) computationally expensive

Hyper-parameters for Ensemble Classifiers

- performance depends on hyper-parameters:
 - (a) `n_estimators`: number of "weak" learners to use
 - (b) `max_features`: maximum number of features for each learner
 - (c) additional hyper-parameters on learners (e.g. `max_depth` of trees in random forest)

A Numerical Dataset

object x_i	Height (H)	Weight (W)	Foot (F)	Label (L)
x_1	5.00	100	6	green
x_2	5.50	150	8	green
x_3	5.33	130	7	green
x_4	5.75	150	9	green
x_5	6.00	180	13	red
x_6	5.92	190	11	red
x_7	5.58	170	12	red
x_8	5.92	165	10	red

- $N = 8$ items
- $M = 3$ (unscaled) attributes


Code for the Dataset

```
import pandas as pd
data = pd.DataFrame(
    {'id': [ 1,2,3,4,5,6,7,8],
     'Label': ['green','green','green','green',
               'red','red','red','red'],
     'Height': [5, 5.5, 5.33, 5.75,
                6.00, 5.92, 5.58, 5.92],
     'Weight': [100, 150, 130, 150,
                180, 190, 170, 165],
     'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
    columns = ['id', 'Height', 'Weight',
               'Foot', 'Label'] )
```

```
ipdb> data
```

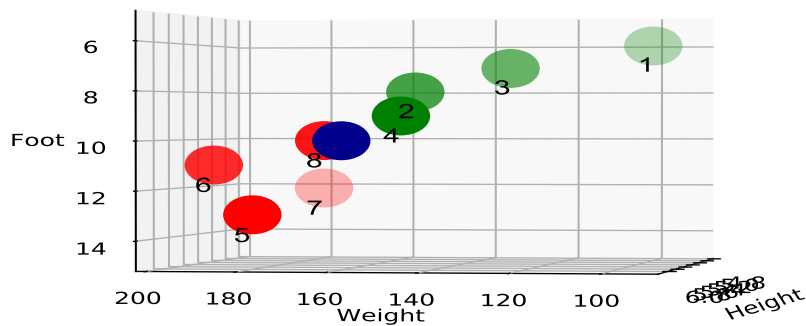
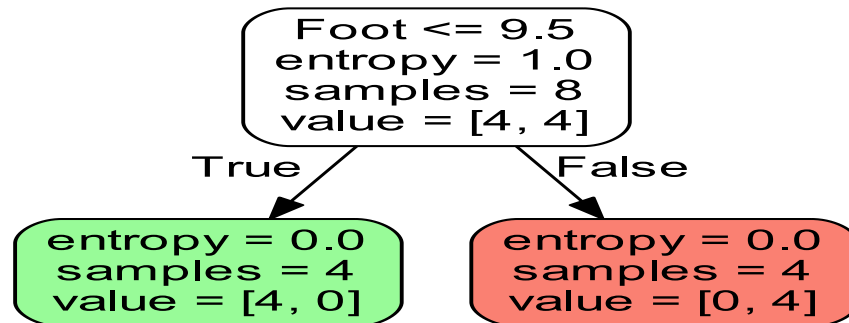
	id	Height	Weight	Foot	Label
0	1	5.00	100	6	green
1	2	5.50	150	8	green
2	3	5.33	130	7	green
3	4	5.75	150	9	green
4	5	6.00	180	13	red
5	6	5.92	190	11	red
6	7	5.58	170	12	red
7	8	5.92	165	10	red

A New Instance



$(H=6, W=160, F=10) \mapsto ?$

Decision Tree



$(H=6, W=160, F=10) \mapsto \text{red}$

Decision Tree in Python

```
import numpy as np
import pandas as pd
from sklearn import tree


data = pd.DataFrame( {'id': [ 1,2,3,4,5,6,7,8],
                      'Label': ['green', 'green', 'green', 'green',
                                'red', 'red', 'red', 'red'],
                      'Height': [5, 5.5, 5.33, 5.75,
                                6.00, 5.92, 5.58, 5.92],
                      'Weight': [100, 150, 130, 150,
                                180, 190, 170, 165],
                      'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
                    columns = ['id', 'Height', 'Weight',
                              'Foot', 'Label'] )

X = data[['Height', 'Weight', 'Foot']].values
Y = data[['Label']].values
clf = tree.DecisionTreeClassifier(criterion = 'entropy')
clf = clf.fit(X,Y)

prediction = clf.predict(np.asmatrix([6, 160, 10]))
```

```
ipdb> prediction[0]
'red'
```


F/W/H Change



id	Height	Weight	Foot	Label
1	5 \mapsto 6	100 \mapsto 170	6 \mapsto 10	green


$(H=6, W=160, F=10) \mapsto ?$

Intermediate Decision



- $(H=6, W=160, F=10) \mapsto \text{green}$
- decide by weight, foot, weight

Decision Tree for F/W/H Change



$(H=6, W=160, F=10) \mapsto \text{green}$

Code for F/W/H Change

```
import numpy as np
import pandas as pd
from sklearn import tree

data = pd.DataFrame( {'id': [ 1,2,3,4,5,6,7,8],
                      'Label': ['green', 'green', 'green', 'green',
                                'red', 'red', 'red', 'red'],
                      'Height': [5, 5.5, 5.33, 5.75,
                                6.00, 5.92, 5.58, 5.92],
                      'Weight': [100, 150, 130, 150,
                                180, 190, 170, 165],
                      'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
                      columns = ['id', 'Height', 'Weight',
                                'Foot', 'Label'] )

data['Foot'].iloc[1] = 10      # change foot from 6 to 10!
data['Weight'].iloc[1] = 170  # weight from 100 to 170
data['Height'].iloc[1] = 6    # height from 5 to 6

X = data[['Height', 'Weight', 'Foot']].values
Y = data[['Label']].values
clf = tree.DecisionTreeClassifier(criterion = 'entropy')
clf = clf.fit(X,Y)
prediction = clf.predict(np.asmatrix([6, 160, 10]))
```

```
ipdb> prediction[0]
```

```
'green'
```

RF in Python

```
import numpy as np
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
data = pd.DataFrame( {'id': [ 1,2,3,4,5,6,7,8],
                      'Label': ['green', 'green', 'green', 'green',
                                'red', 'red', 'red', 'red'],
                      'Height': [5, 5.5, 5.33, 5.75,
                                6.00, 5.92, 5.58, 5.92],
                      'Weight': [100, 150, 130, 150,
                                180, 190, 170, 165],
                      'Foot': [6, 8, 7, 9, 13, 11, 12, 10]},
                    columns = ['id', 'Height', 'Weight',
                              'Foot', 'Label'] )
data['Foot'].iloc[1] = 10; # foot from 6 to 10
data['Weight'].iloc[1] = 170 # weight from 100 to 170
data['Height'].iloc[1] = 6 # height from 5 to 6
X = data[['Height', 'Weight', 'Foot']].values
Y = data[['Label']].values
class_labels_dict = {'green': 1, 'red': 0}
label_color_dict = {1: 'green', 0: 'red'}
data['class_labels'] = data['Label'].map(class_labels_dict)
model = RandomForestClassifier(n_estimators=5, max_depth=3,
                              criterion='entropy')

model.fit(X, Y)
test_instance = np.asmatrix([6, 160, 10])
rf_label = int(model.predict(test_instance)[0])
rf_color = label_color_dict[rf_label]
```

```
ipdb> rf_color
```

```
'red'
```


Python Code

```
for i in range(5):
    estimator = model.estimators_[i]
    rf_label = int(estimator.predict(test_instance)[0])
    rf_color = label_color_dict[label_prediction]
    print('estimator:', i, ' label = ', rf_label, ' color: ',
          rf_color_color)
```

ipdb>

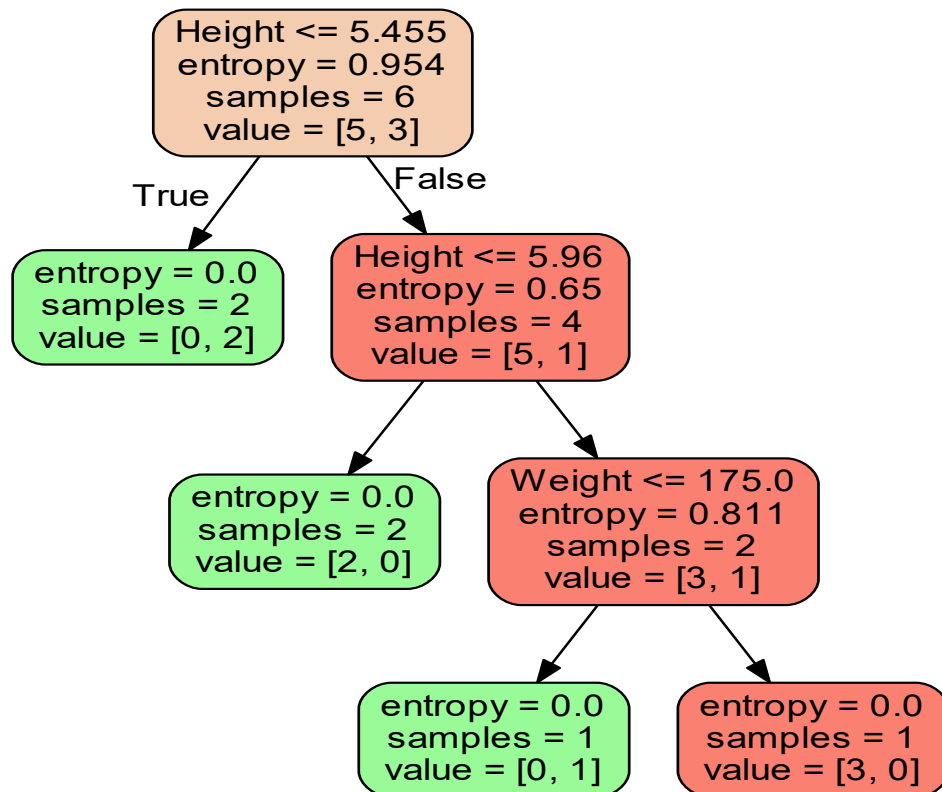
```
estimator: 0  label = 0  color:  red
estimator: 1  label = 0  color:  green
estimator: 2  label = 1  color:  green
estimator: 3  label = 1  color:  red
estimator: 4  label = 0  color:  red
```

RF Estimator 0



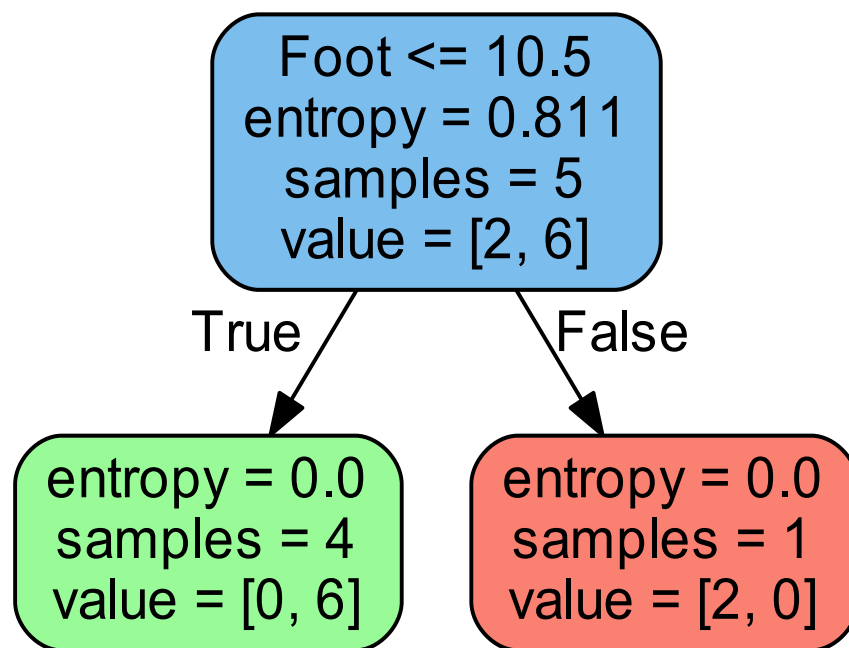
$(H=6, W=160, F=10) \mapsto \text{red}$

RF Estimator 1



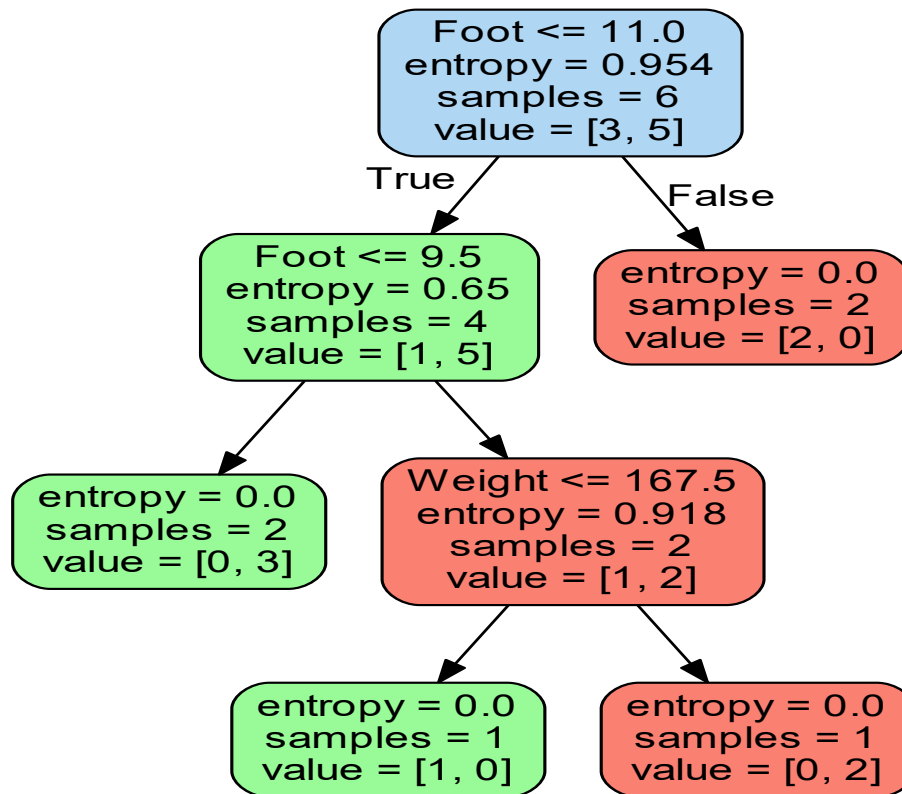
$(H=6, W=160, F=10) \mapsto \text{green}$

RF Estimator 2



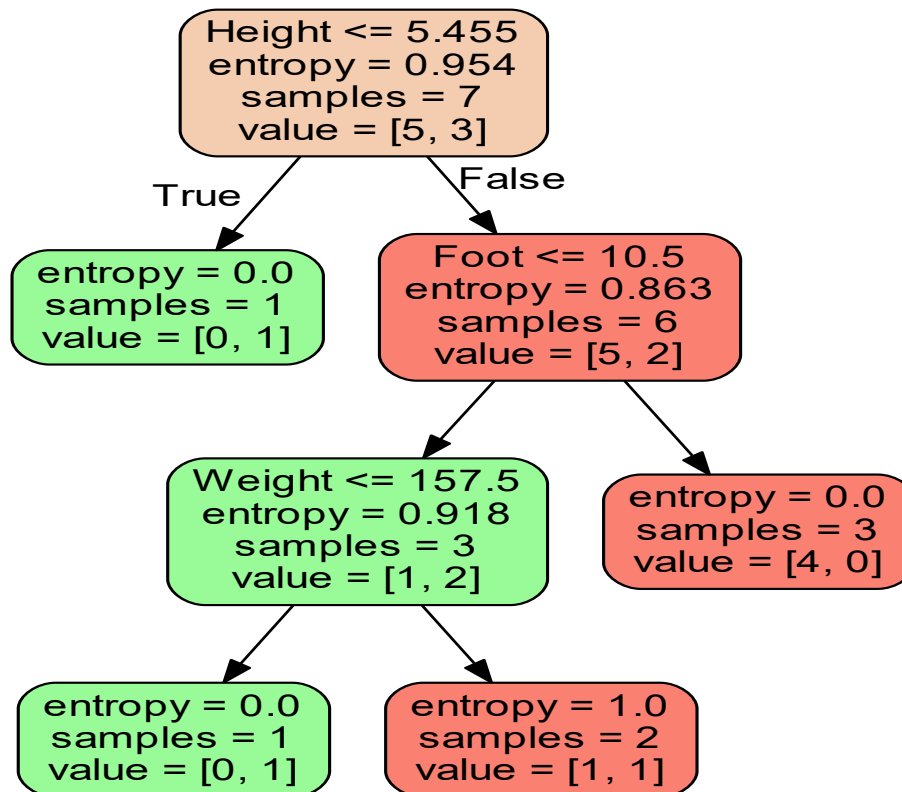
$(H=6, W=160, F=10) \mapsto \text{green}$

RF Estimator 3



$(H=6, W=160, F=10) \mapsto \text{red}$

RF Estimator: 4



$(H=6, W=160, F=10) \mapsto \text{red}$

Random Forest: IRIS

```
import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

url = r'https://archive.ics.uci.edu/ml/' + \
      r'machine-learning-databases/iris/iris.data'

iris_feature_names = ['sepal-length', 'sepal-width',
                      'petal-length', 'petal-width']
data = pd.read_csv(url, names=['sepal-length', 'sepal-width',
                              'petal-length', 'petal-width', 'Class'])

class_labels = ['Iris-versicolor', 'Iris-virginica']
data = data[data['Class'].isin(class_labels)]

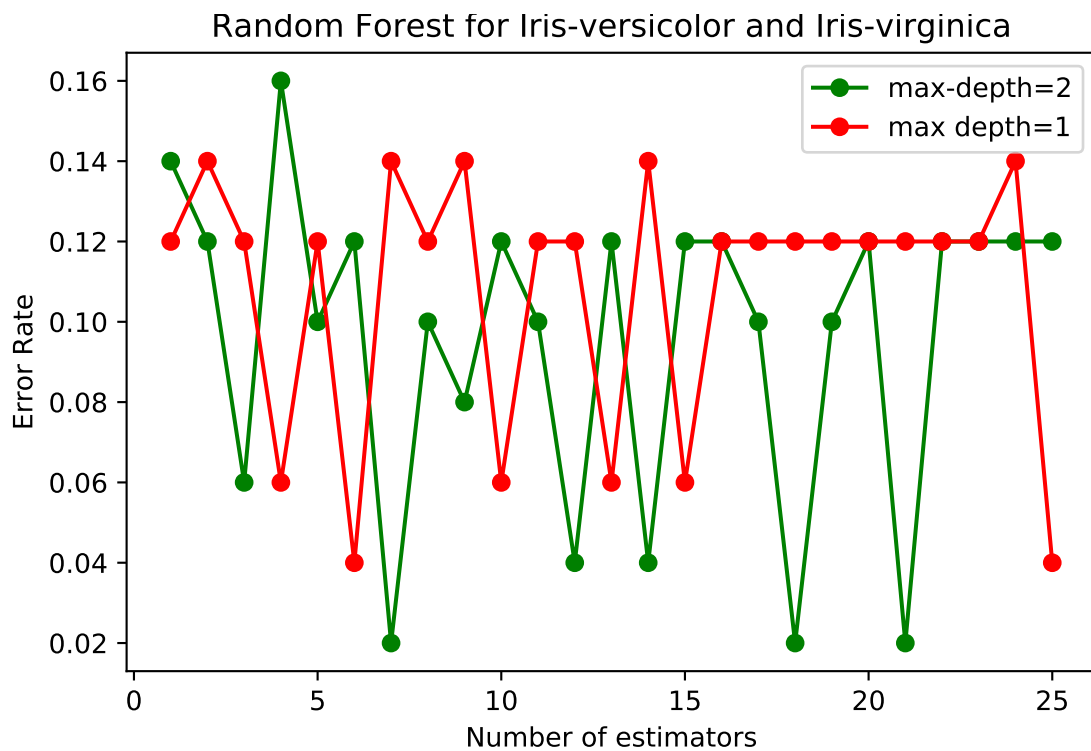
X = data[iris_feature_names].values
le = LabelEncoder()
Y = le.fit_transform(data['Class'].values)

X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
                                                    test_size=0.5, random_state=3)

model = RandomForestClassifier(n_estimators=25, max_depth=5,
                              criterion='entropy')
model.fit(X_train, Y_train)
prediction = model.predict(X_test)
error_rate = np.mean(prediction != Y_test)

ipdb> error_rate
0.1
```

Impact of Depth and Number of Estimators



AdaBoost Method

- choose a base classifier
- use k such classifiers C_1, \dots, C_k (“weak” learners)
- each C_i does a prediction on a subset X_i of data
- assign weight w_i to C_i based on its accuracy
- use weighted average of predictions

AdaBoost: IRIS

```
import pandas as pd
import numpy as np
from sklearn.svm import SVC # use SVM as base
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

url = r'https://archive.ics.uci.edu/ml/' + \
      r'machine-learning-databases/iris/iris.data'

iris_feature_names = ['sepal-length', 'sepal-width',
                      'petal-length', 'petal-width']
data = pd.read_csv(url, names=['sepal-length', 'sepal-width',
                              'petal-length', 'petal-width', 'Class'])

class_labels = ['Iris-versicolor', 'Iris-virginica']
data = data[data['Class'].isin(class_labels)]

X = data[iris_feature_names].values
le = LabelEncoder()
Y = le.fit_transform(data['Class'].values)


X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
                                                    test_size=0.5, random_state=3)

svc=SVC(probability=True, kernel='linear') # use as base learner
model = AdaBoostClassifier(n_estimators=5, base_estimator=svc,
                           learning_rate = 0.5)

model.fit(X_train, Y_train)
prediction = model.predict(X_test)
error_rate = np.mean(prediction != Y_test)

ipdb> error_rate
0.02
```

Impact of Learning Rate and Learners



Concepts Check:

- (a) ensemble learning
- (b) bagging
- (c) advantages and disadvantages
- (d) hyperparameters (estimators, max features, depth)
- (e) Random Forest classification
- (f) AdaBoost classification